

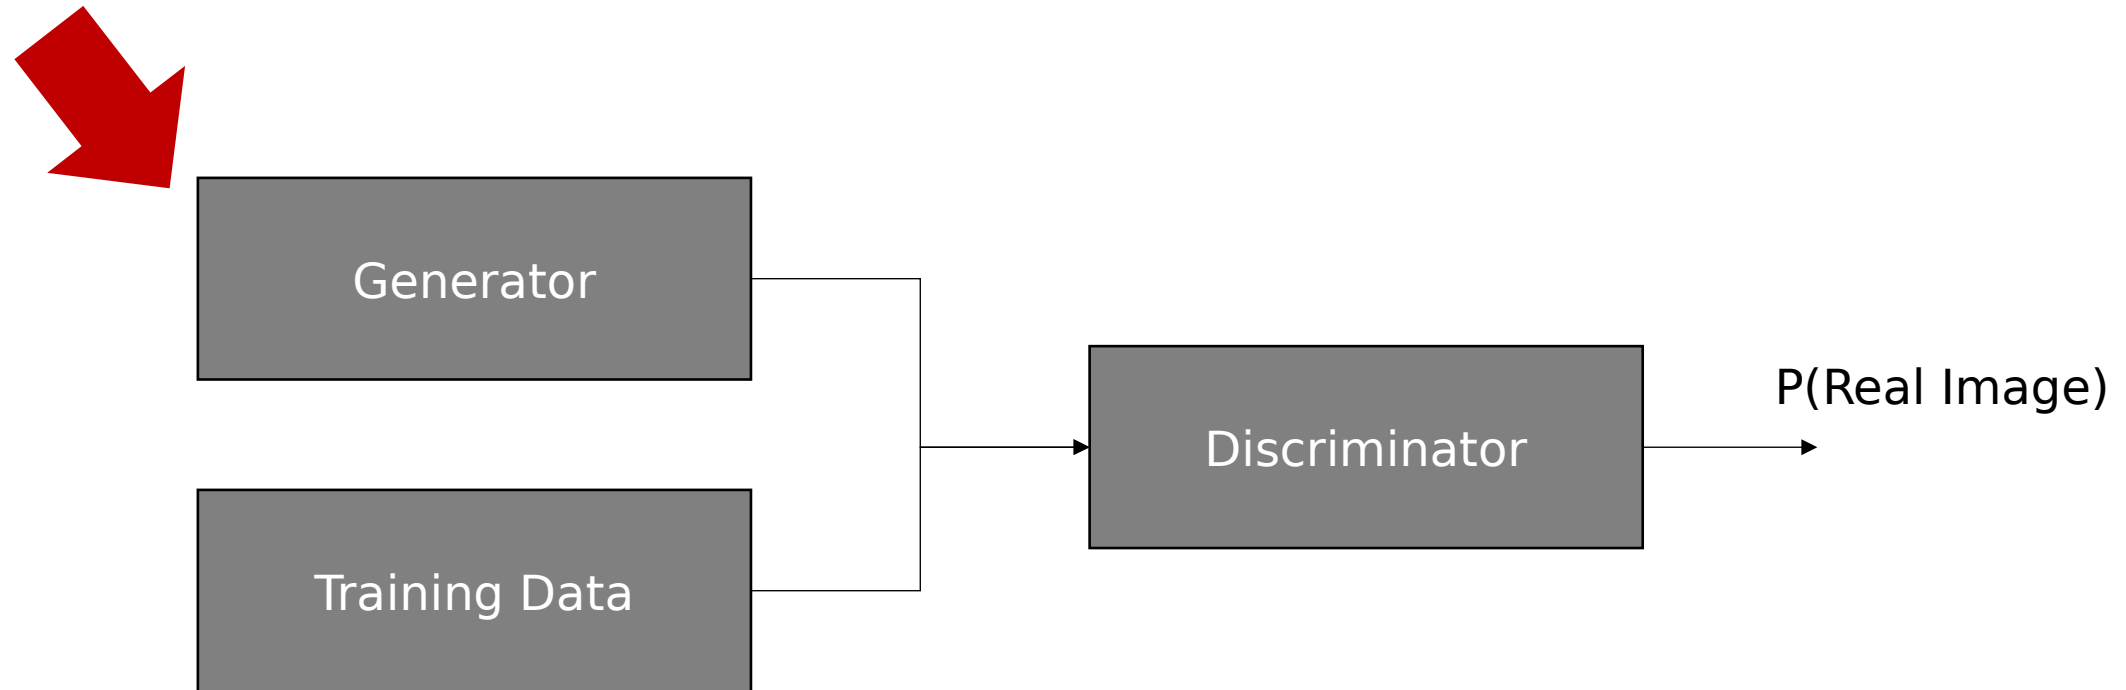
# A Style-Based Generator Architecture for Generative Adversarial Networks

Tero Karras, Samuli Laine, Timo Aila  
NVIDIA

Presenter: Diego Cantor, PhD

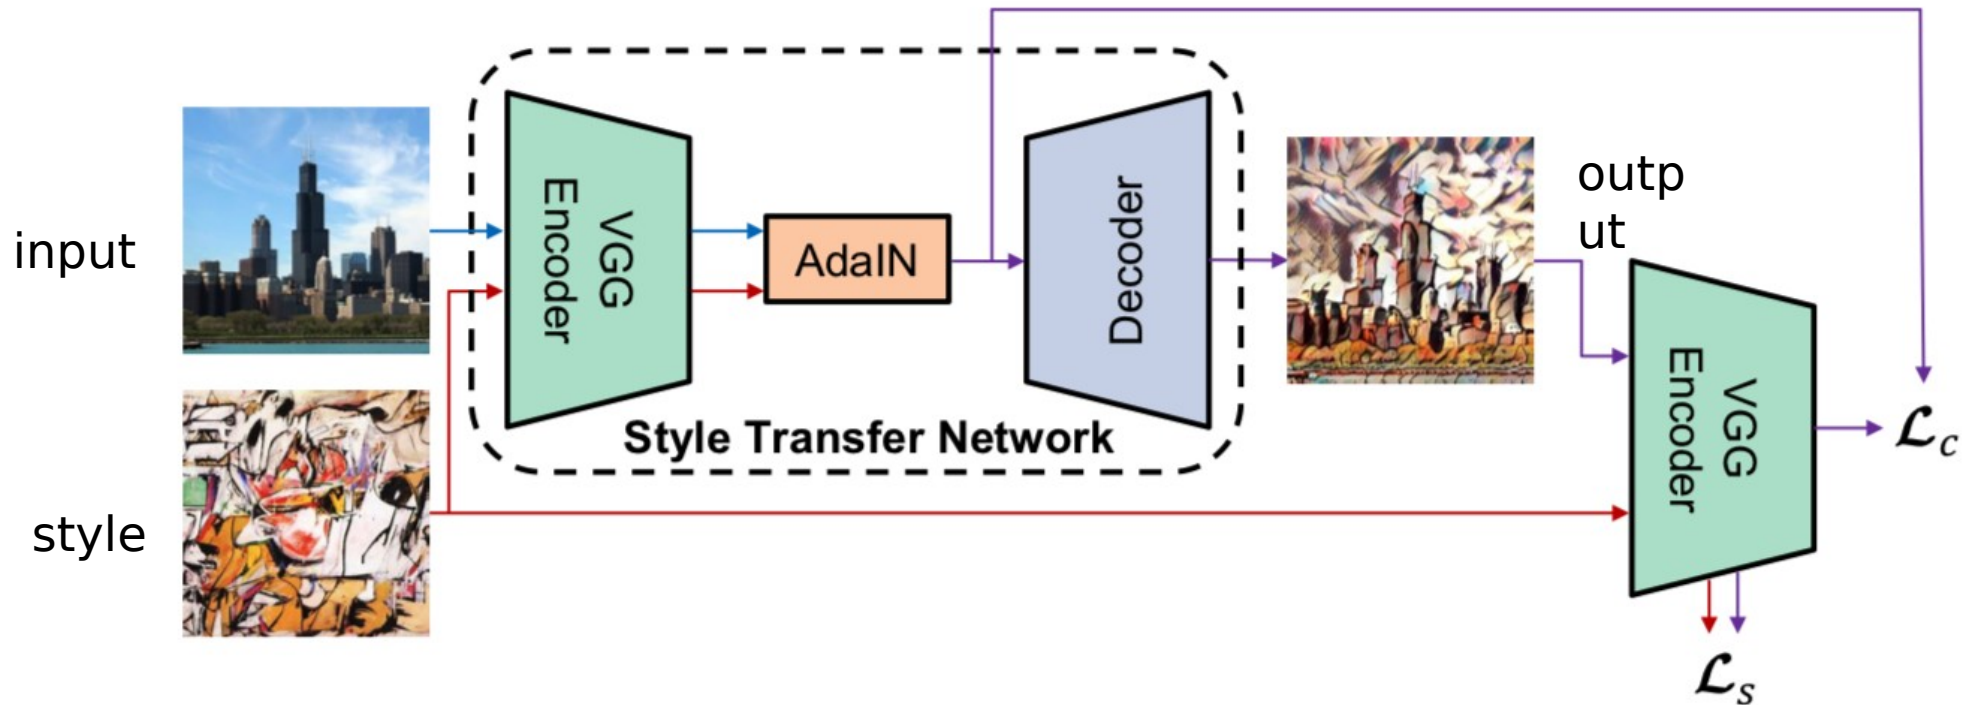
Facilitators: Michael Vertolli and David McDonald

Despite improvement in image quality synthesis,  
GAN generators operate as black boxes



Understanding of image synthesis  
is poor

# This work proposes a model for the generator that is inspired by **style transfer networks**



Arbitrary Style Transfer in Real-time with Adaptive Instance Normalization, Huang and Belongie, 2017


# Everything started with the usage of batch normalization to improve training

$$\text{BN}(x) = \gamma \left( \frac{x - \mu(x)}{\sigma(x)} \right) + \beta$$



$$\text{IN}(x) = \gamma \left( \frac{x - \mu(x)}{\sigma(x)} \right) + \beta$$

Gamma and Beta  
are learned from  
data



Arbitrary Style Transfer in Real-time with Adaptive Instance Normalization, Huang and Belongie

# Instance normalization improves style-transfer loss when compared to other approaches



Content




Style



StyleNet BN



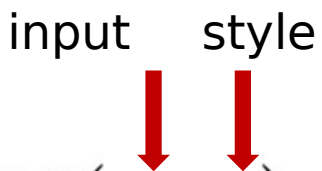
StyleNet IN (ours)



Improved Texture Networks: Maximizing Quality and Diversity in Feed-forward Stylization and Texture Synthesis, Ulyanov and Vedaldi, CVPR, 2017

Adaptive Instance Normalization simply scales the normalized input with style spatial statistics. This has profound implications.


input style


$$\text{AdaIN}(x, y) = \sigma(y) \left( \frac{x - \mu(x)}{\sigma(x)} \right) + \mu(y)$$

Style statistics are not learnable. So  
AdaIN has no learnable parameters

$$\text{BN}(x) = \gamma \left( \frac{x - \mu(x)}{\sigma(x)} \right) + \beta \qquad \text{IN}(x) = \gamma \left( \frac{x - \mu(x)}{\sigma(x)} \right) + \beta$$







$$\text{AdaIN}(x, y) = \sigma(y) \left( \frac{x - \mu(x)}{\sigma(x)} \right) + \mu(y)$$



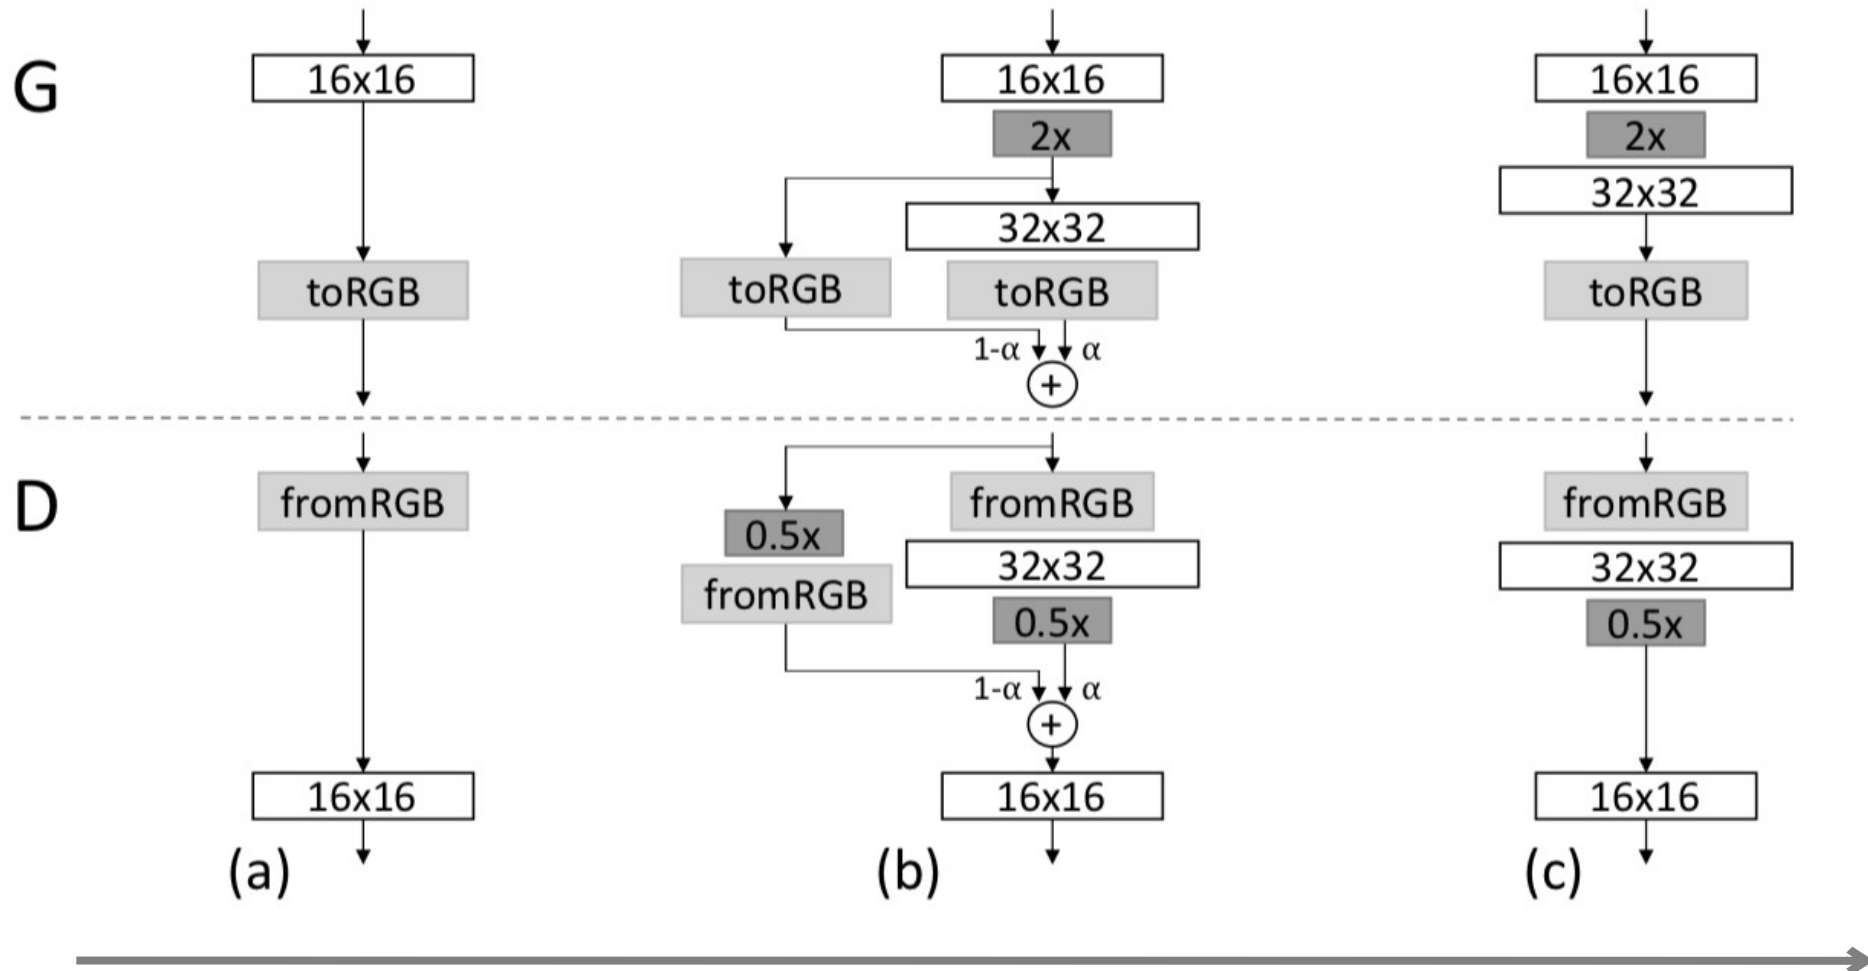
# The baseline configuration is the progressive GAN setup (same research group at NVIDIA)

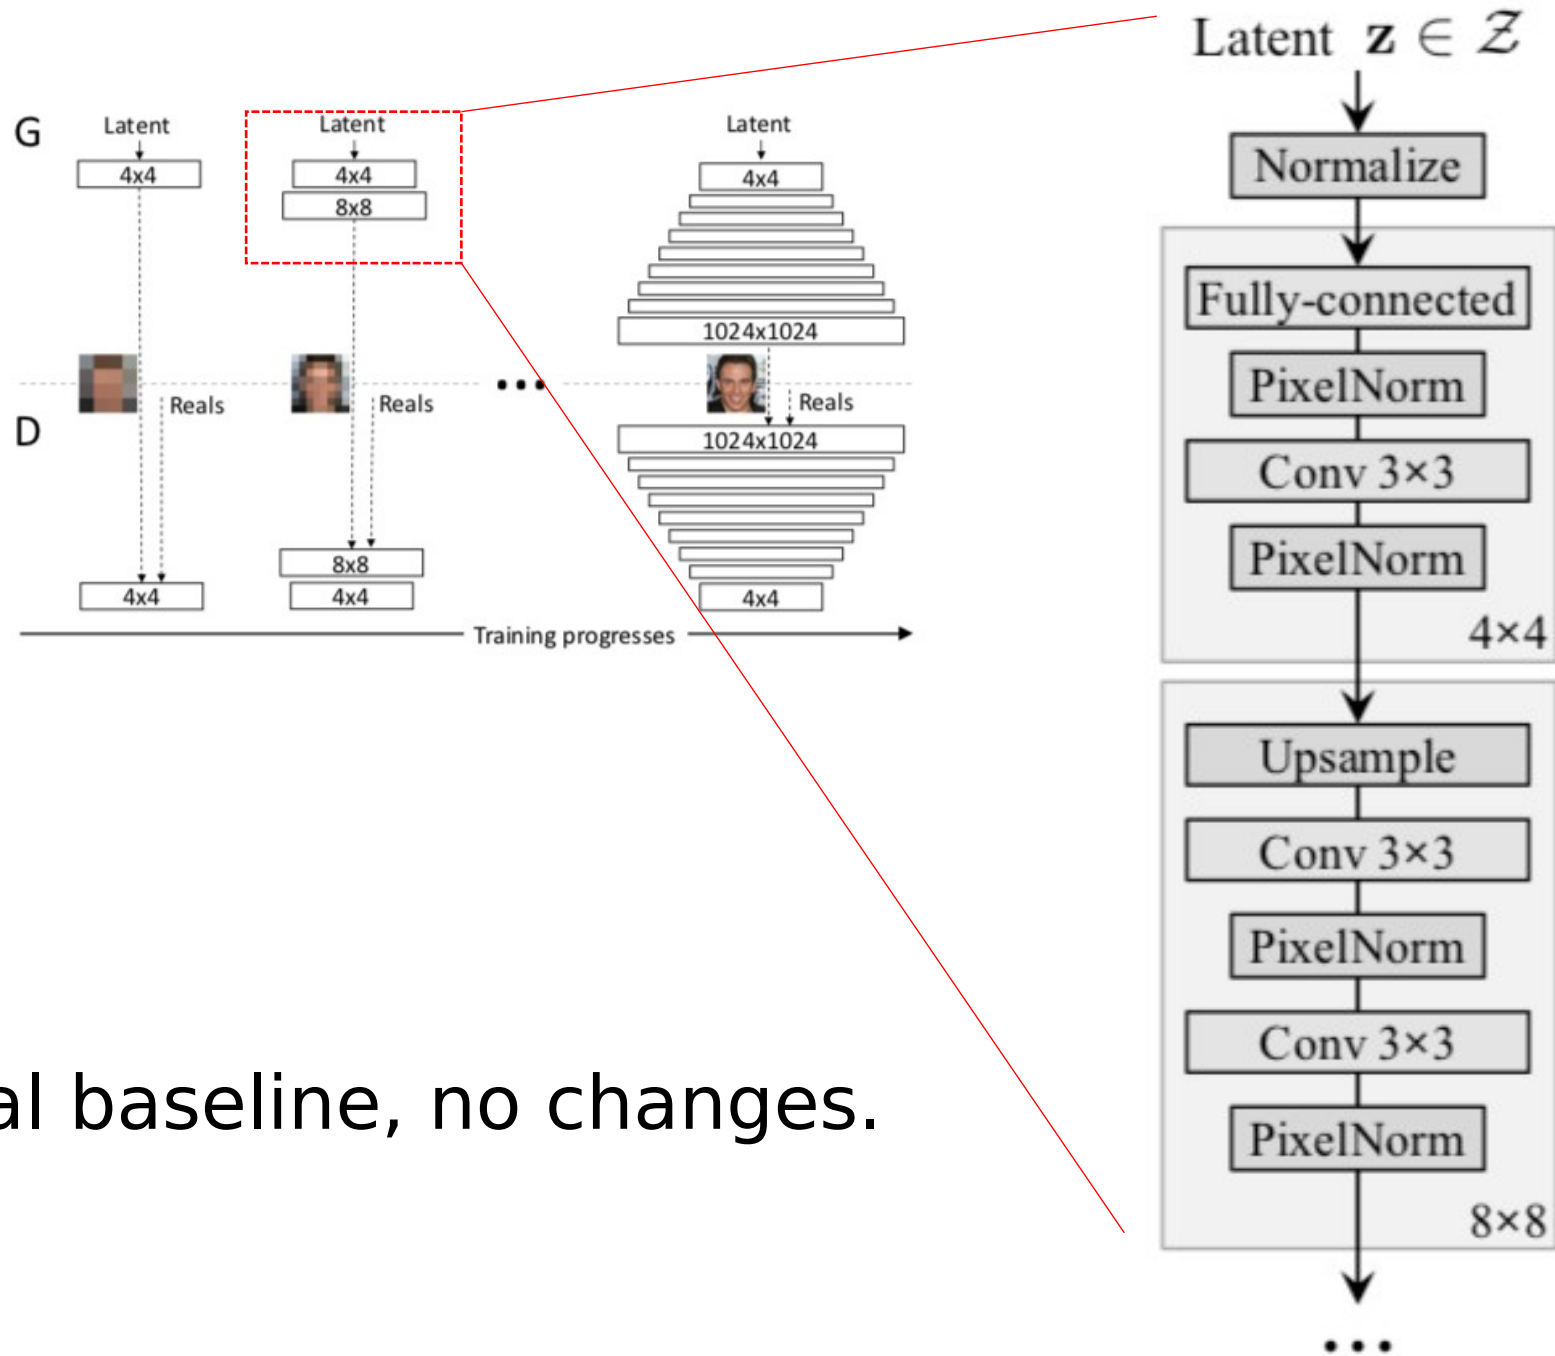


Progressive growing of GANs for improved quality, stability and variation, Karras et al., ICLR 2018



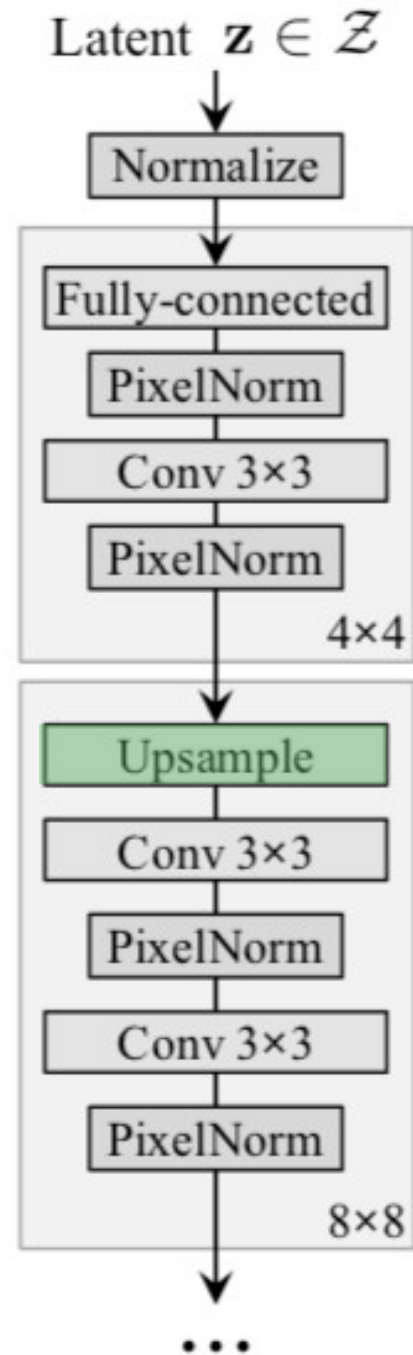
# Smooth transition into higher-res layers using bilinear interpolation





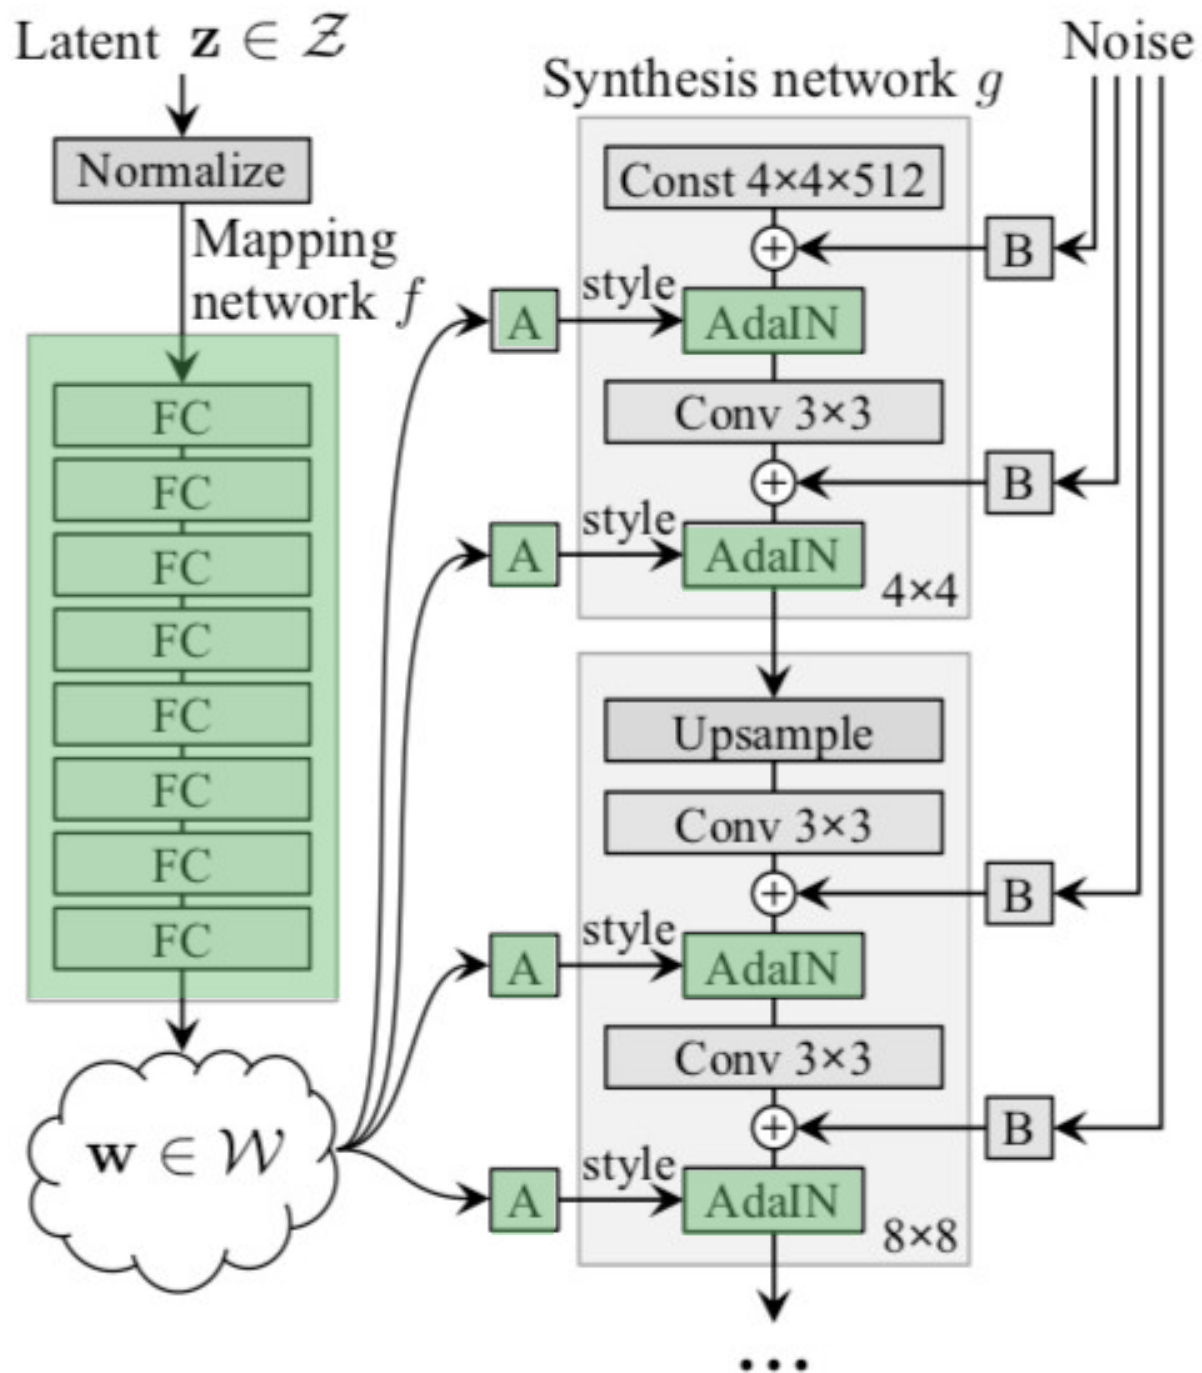
Step A

Original baseline, no changes.



## Step B

- Replace nearest neighbor with bilinear upsampling
- Replace pooling with bilinear downsampling (in the discriminator)

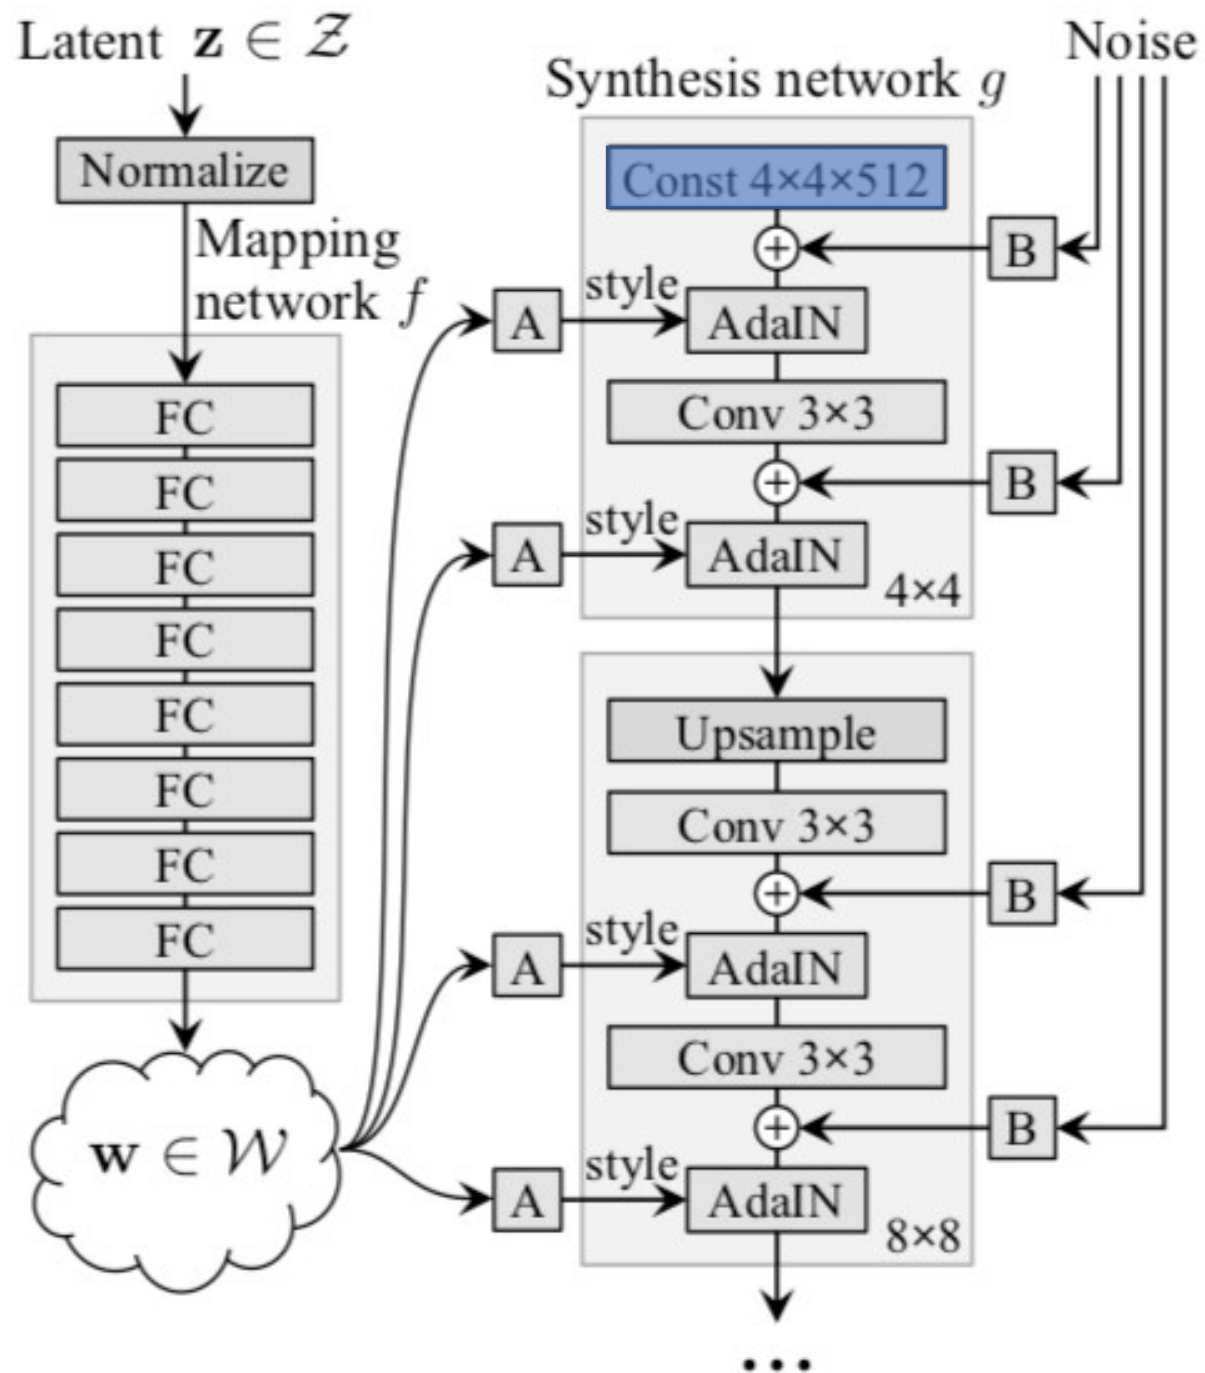


Step C

Add mapping network and styles.

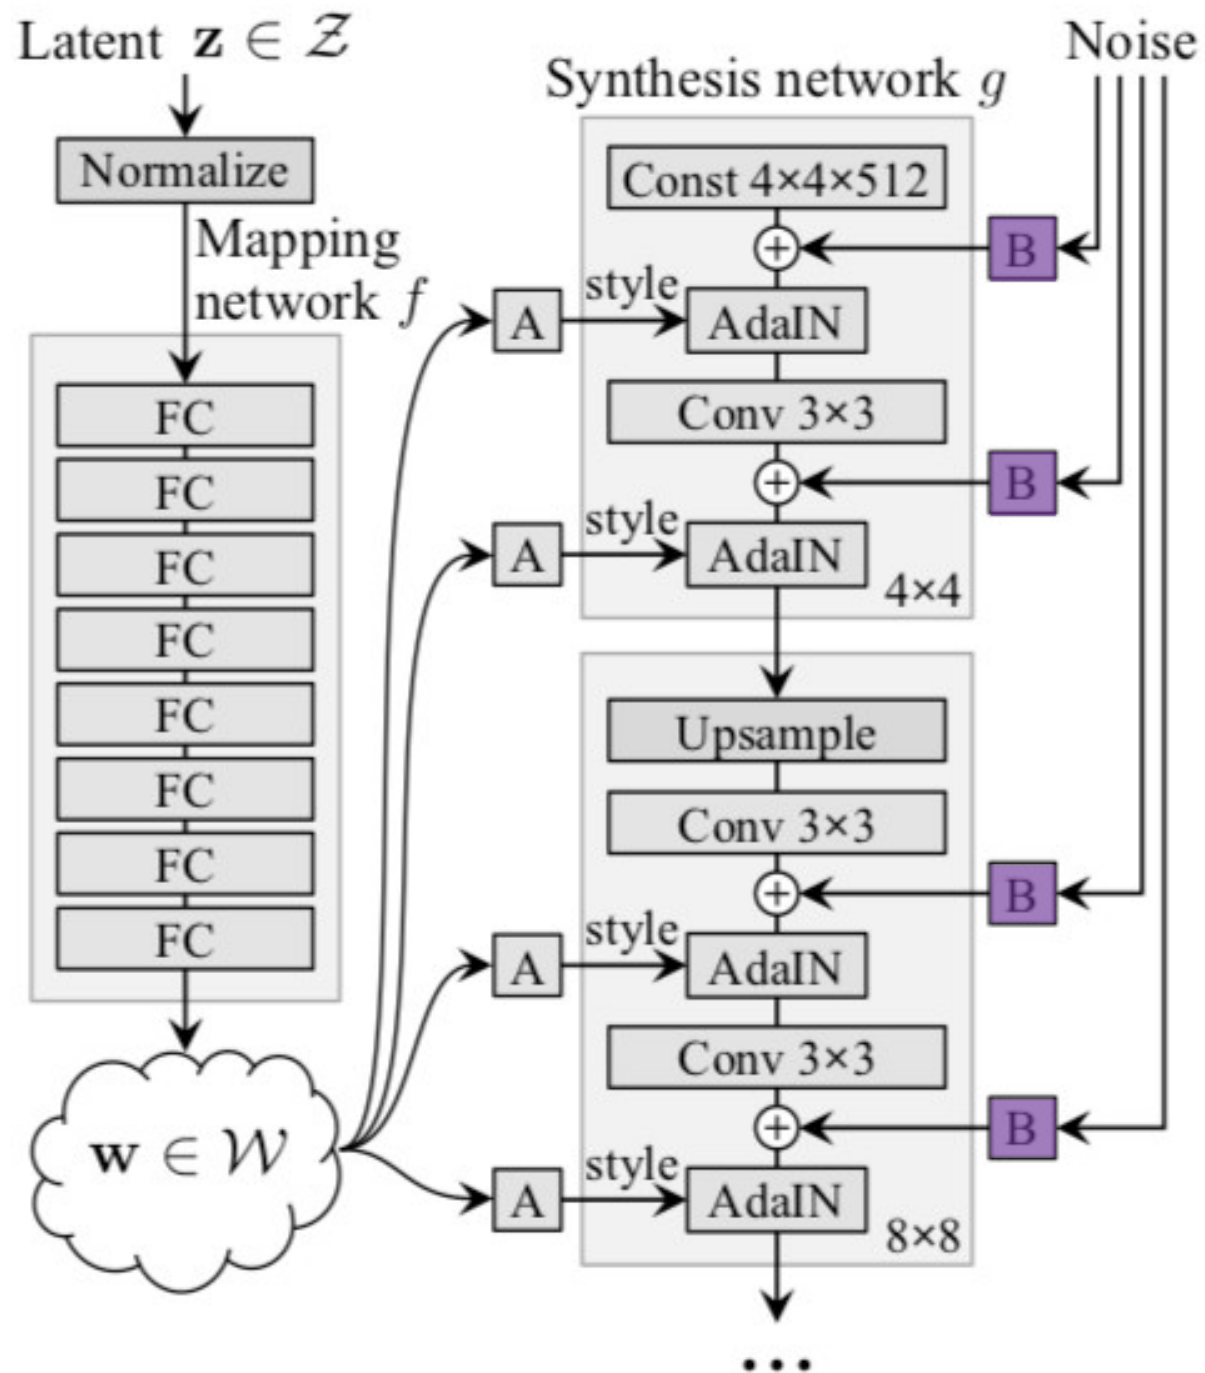
**Styles** are generated from  $\mathbf{W}$  and used in AdaIN operations





Step D

remove traditional  
input

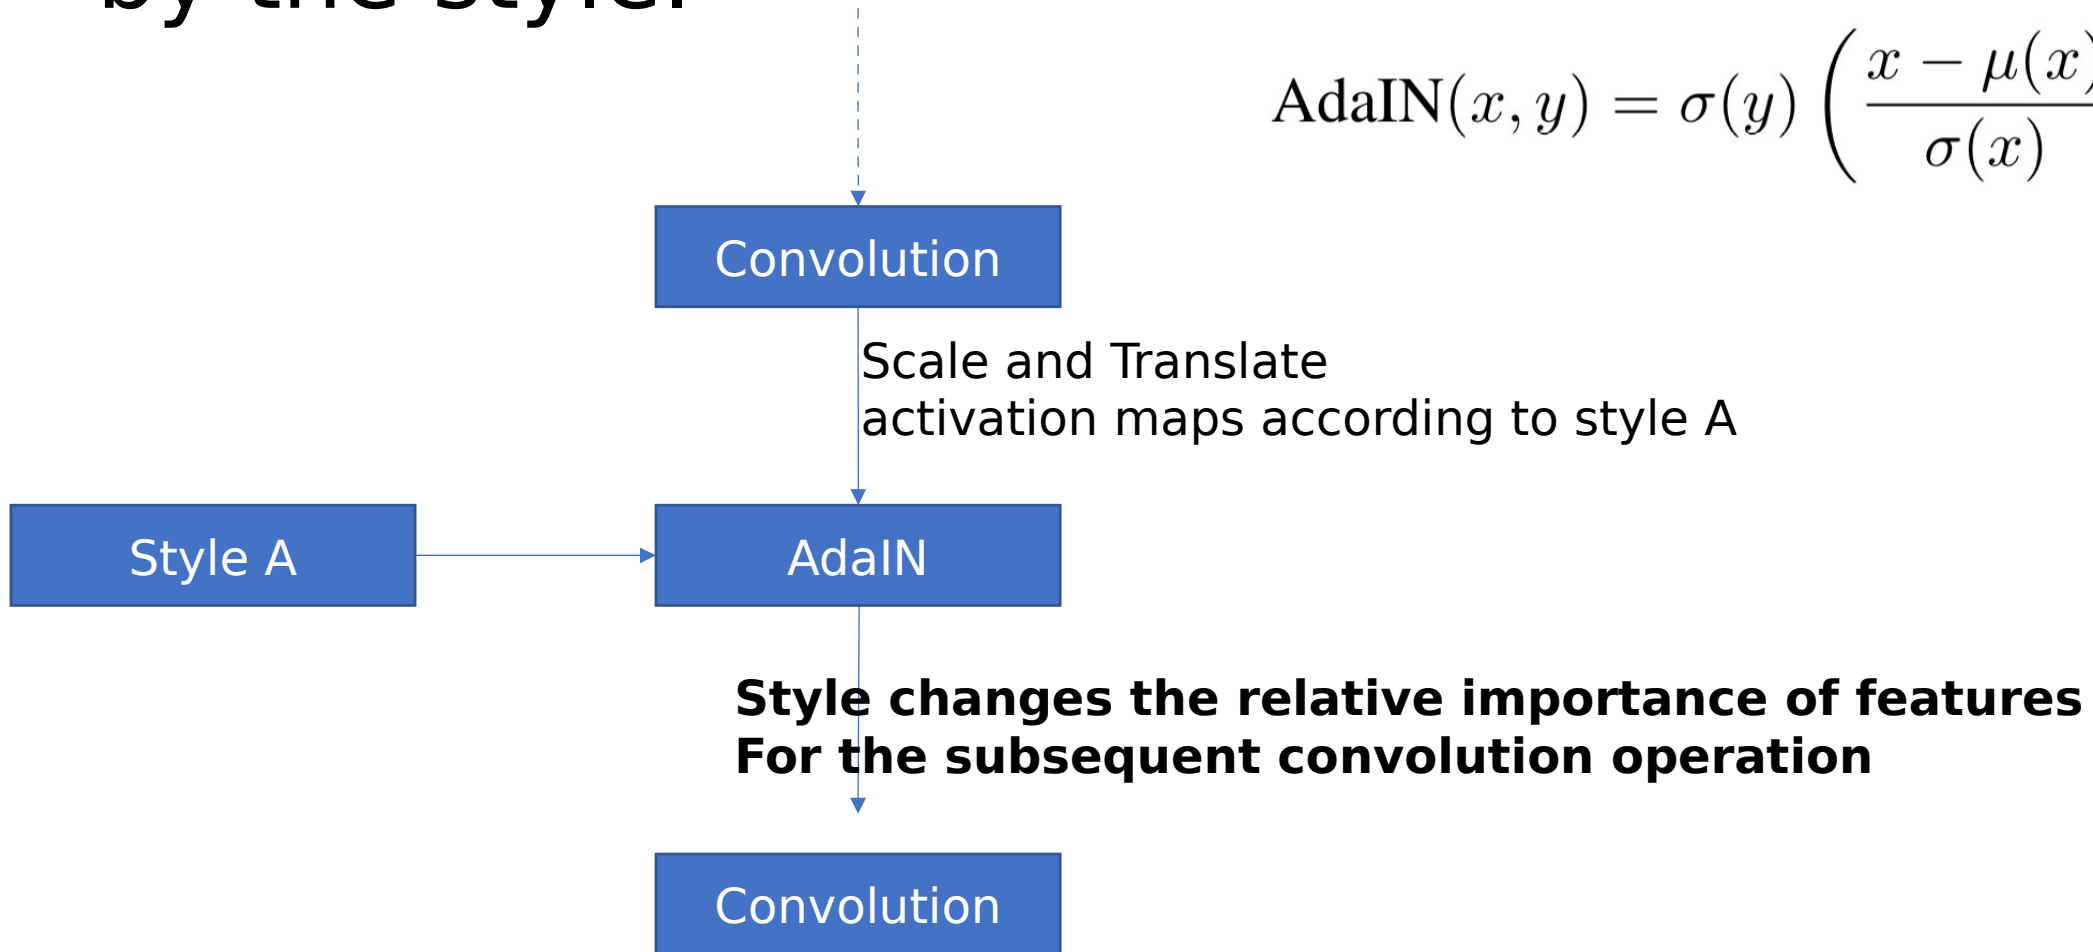


Step E



Add noise inputs  
(enables  
generating  
stochastic detail)

This is the key: AdaIN operation **affects the relative importance of features at every scale**. How much? This is determined by the style.

$$\text{AdaIN}(x, y) = \sigma(y) \left( \frac{x - \mu(x)}{\sigma(x)} \right) + \mu(y)$$



Style affects the entire image but noise is added per pixel. The network learns to use it to control **stochastic variation**.



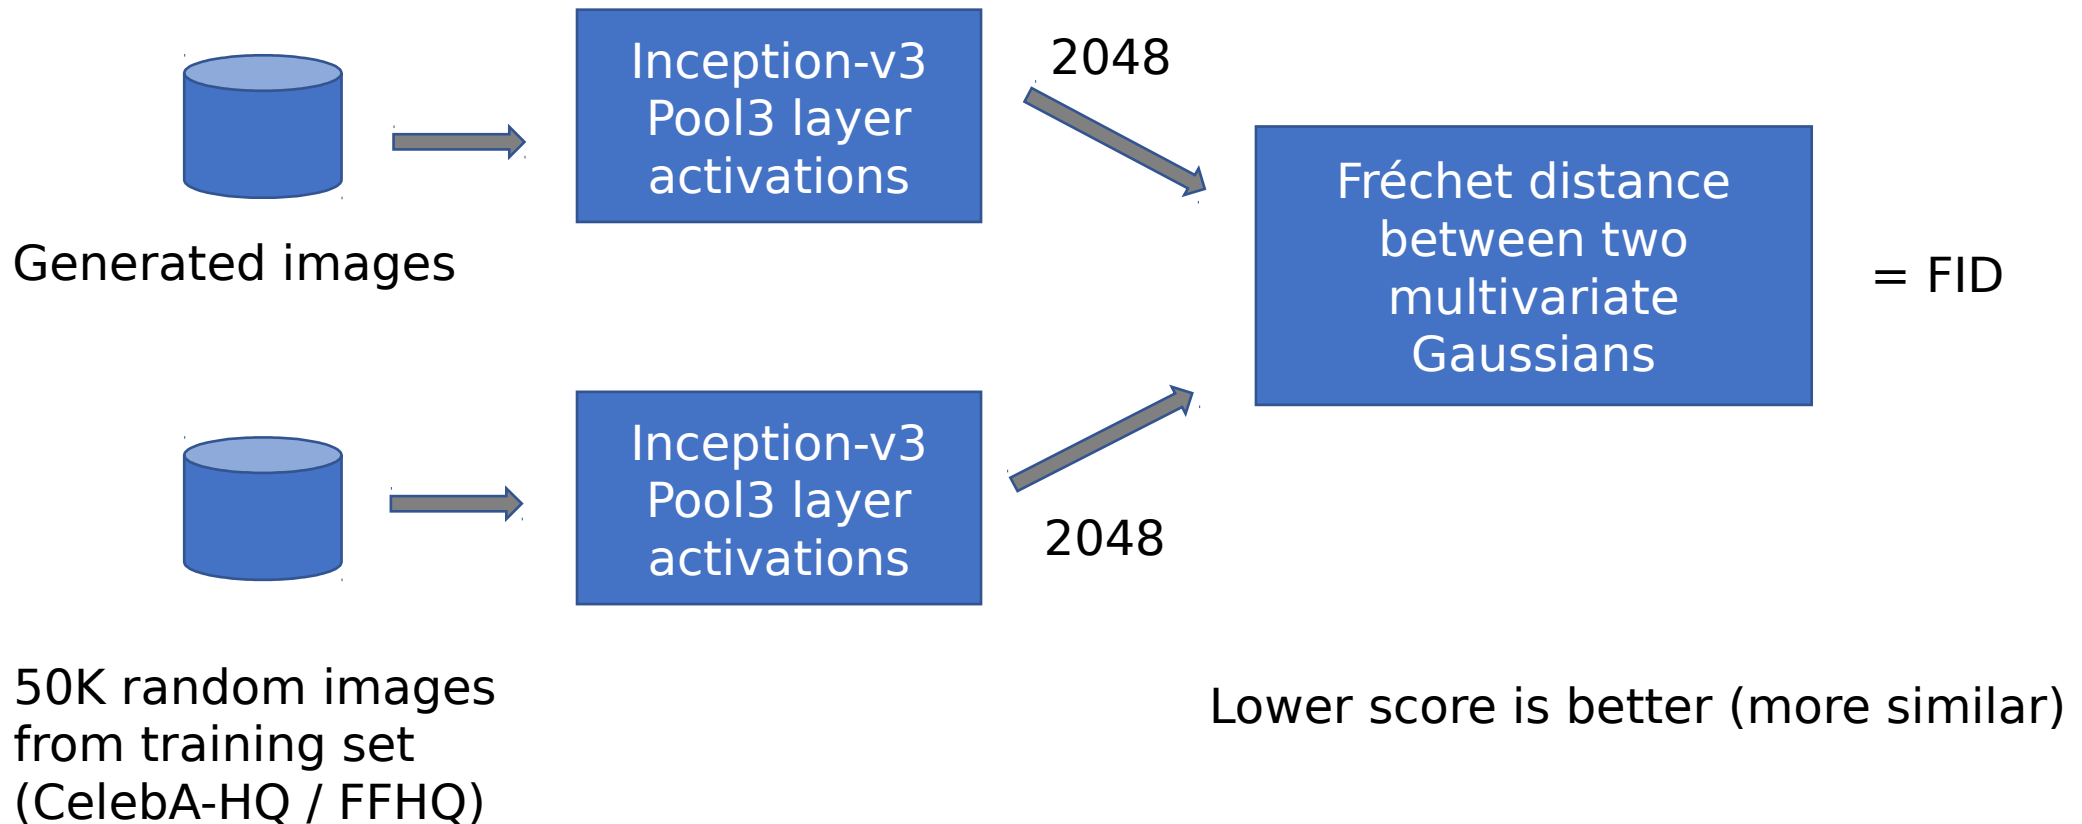


2 min break

# Results



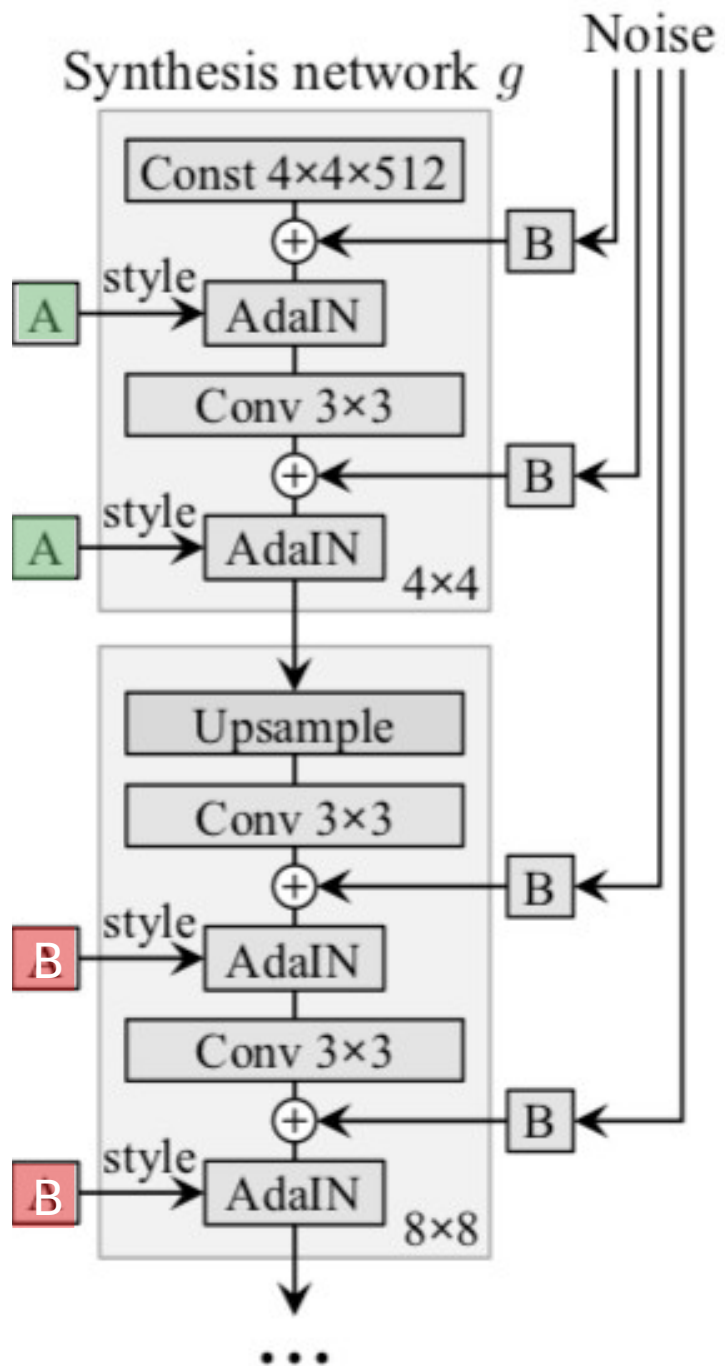
This group used the **Fréchet inception distance** (FID) to measure the quality of generated images




Results: quality of the generated images. Lower FID is better

Method	CelebA-HQ	FFHQ
A Baseline Progressive GAN	7.79	8.04
B + Tuning (incl. bilinear up/down)	6.11	5.25
C + Add mapping and styles	5.34	4.85
D + Remove traditional input	5.07	4.88
E + Add noise inputs	<b>5.06</b>	4.42
F + Mixing regularization	5.17	<b>4.40</b>





# Mixing styles during image synthesis



Mixing styles during image synthesis. Coarse styles such as pose, face shape and glasses are copied.





# Middle styles copied: hair style, facial features but not pose or glasses






# Copying only fine resolution style such as colour scheme





# Style-based generator architecture



# Major contributions

1. Significant improvement over traditional GAN generators architecture
2. Separation of high-level attributes from stochastic effects
3. Does not generate new images **from scratch** but rather through a smart combination of styles that are embedded in sample images (latent codes)