
TDLs Code Review

Transformer

(Attention Is All You Need)

Lead by **Xiyang Chen**

Facilitated by **Felipe Pérez & Ehsan Amjadian**

Outline

- Motivation
- PyTorch vs. Tensorflow
- Overview
- Individual parts
- Training & experiments
- Discussions

We will be mainly focusing on code.

For more detailed explanation, refer to a [previous TDLS session on Transformer](#)
presented by **Joseph Palermo**.

Why code review

Code review enables us to...

- ...look at details that are glossed over
 - Details only to be found on prior work
 - Practical constraints
 - Memory
 - Wall-clock time (a.k.a. paper deadline)
 - Source code availability
- ...use code to aid understanding
- ...use small experiments to test our assumptions
 - Poke, observe, believe
- Be lazy

Code reviews are new and experimental.
Suggestions are welcome!

Why Transformer

- In a way, it's attention to the extreme
- Achieves SotA's in sequence-related tasks
 - BERT
 - TransferTransfo (convo dialog generation)
 - Transformer-XL
- Foundation for many pioneering works
 - Image Transformer
 - Self-attention CycleGAN
 - AlphaStar
 - Cited by over 1,000 works as of early Feb
- Faster, more scalable, more interpretable
 - Unlike RNN, training can be completely parallelized across sequence timesteps

Pytorch vs. Tensorflow

- The official Transformer implementation is in Tensorflow
- Many people prefer PyTorch
- Which framework then?



Andrej Karpathy ✓

@karpathy

Follow



I've been using PyTorch a few months now
and I've never felt better. I have more energy.
My skin is clearer. My eye sight has
improved.

TF vs. PyTorch: practical pros and cons as of Feb 2019

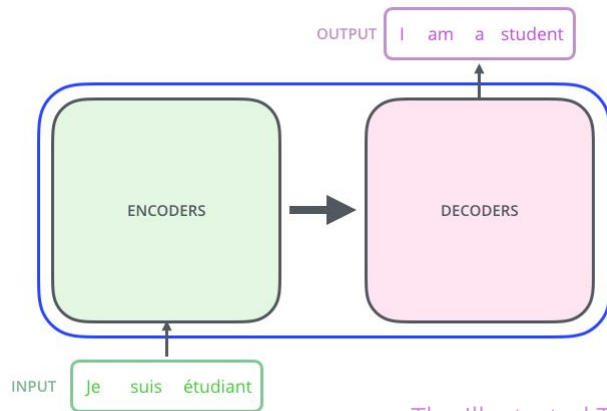
- PyTorch is a very popular choice among researchers
 - Intuitive and flexible
 - Easier to work with CUDA
- TF is production friendly
 - TF Serving gives you both TCP & RESTful APIs
 - TF has more support on most popular cloud platforms (GCP, AWS, etc) in terms of code examples and guides
- TF spans more platforms and types of devices
 - TPU, GPU, CPU, Mobile (TF Lite), Browser/Node.js (TF.js), Edge TPU
- TF's Static Graph mode boosts performance, but is cumbersome to code with, especially for rapid prototyping and experimentation
- TF Eager comes to the rescue
 - API is similar to that of PyTorch and MXNet
 - Can use AutoGraph in Eager Mode to generate graph code
 - Will become the default mode in TF 2.0
 - However, beware: TF Eager is still new. A lot of existing TF code is not compatible with Eager Mode yet

TF vs. PyTorch

- With all considered, we will base our review on the [The Annotated Transformer](#), a PyTorch implementation by Harvard NLP.
- There are other implementations that may be more suitable for your purpose.
 - Check out the reference slide.

Overview

- At high level, an encoder-decoder architecture
- $N = 6$
- Input size: 512
- Output size: 512
- Output size for most layers: 512



[The Illustrated Transformer](#)

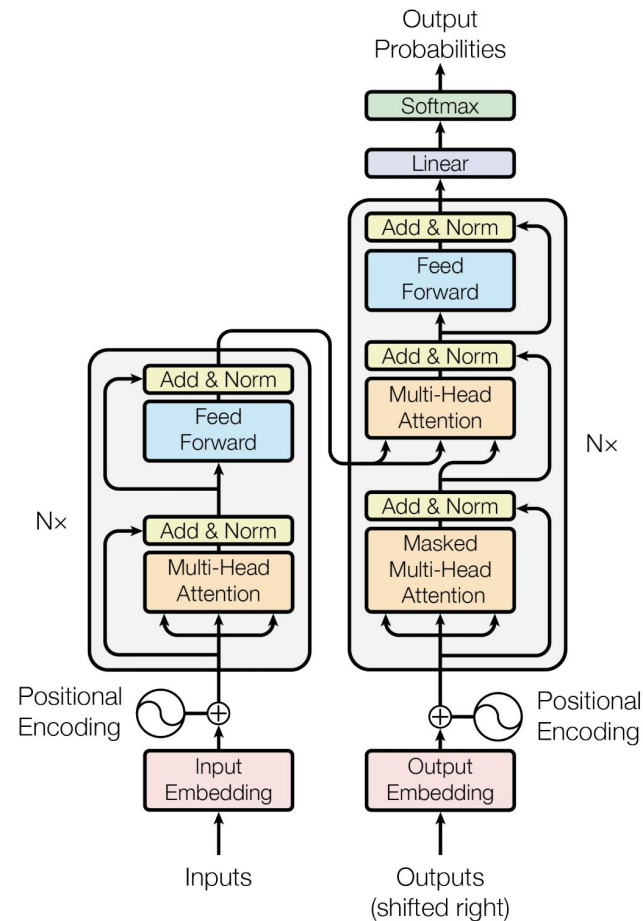
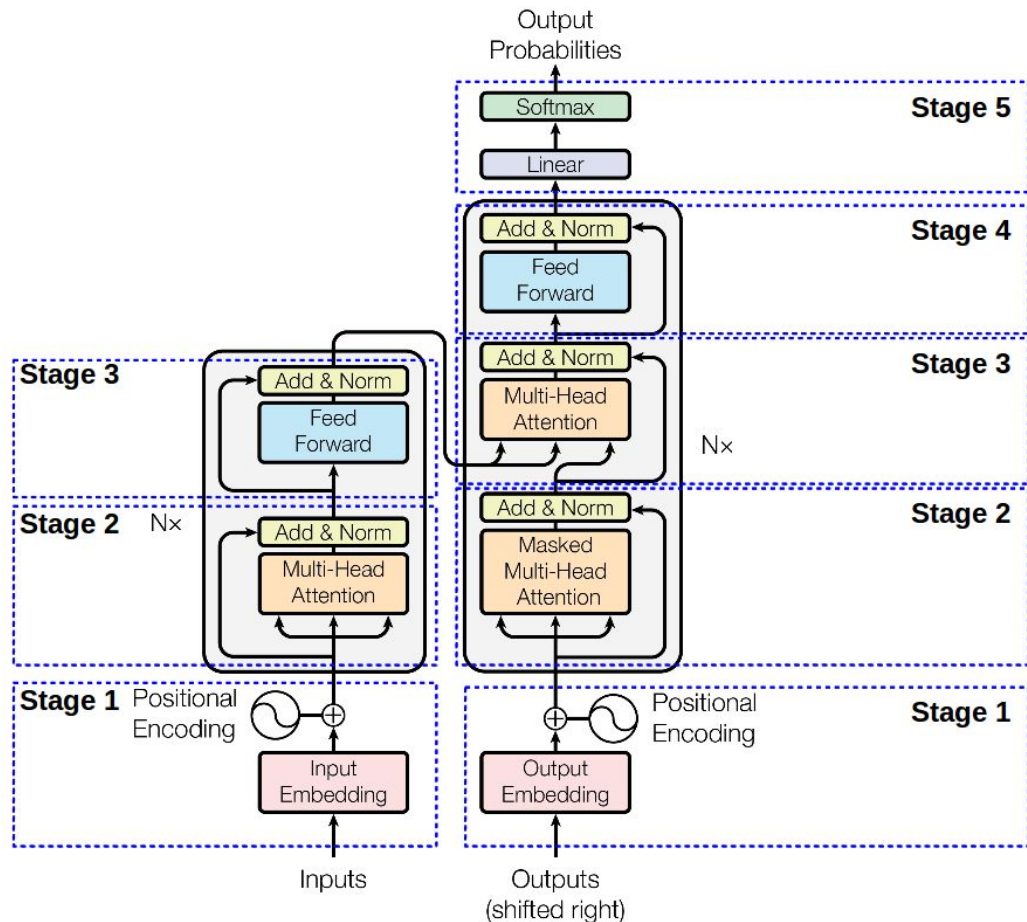


Figure 1: The Transformer - model architecture.

Overview

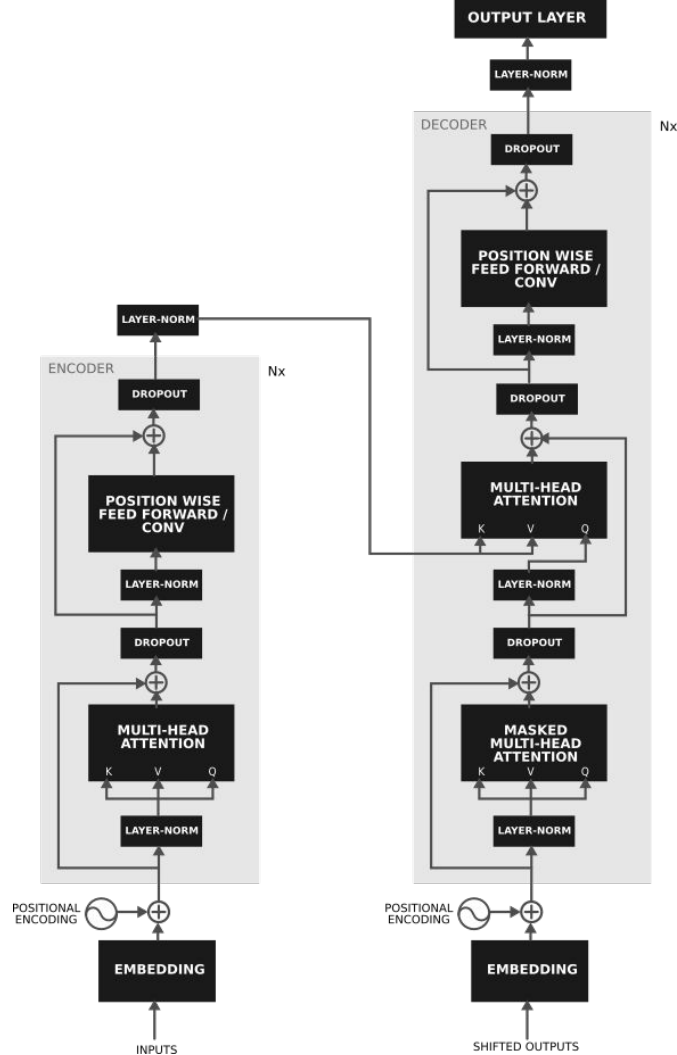
- Source data \Rightarrow encoder
- Target data \Rightarrow decoder
 - Target data is required for learning the context (e.g. what words have been translated so far)
 - Therefore, masked
- Output is compared against target data
 - Loss is $\text{KL_Div}(x, \text{Targ})$



[Michał Chromiak's blog post on Transformer](#)

Overview

- Two unique parts:
 - The Multiheaded Attention layer
 - Positionwise Feed-Forward layer
- Other parts:
 - Positional Encodings
 - Masks
 - Embeddings
- Loss & training
 - Single GPU
 - Multiple GPU



PyTorch Preliminaries

- Modules inherit from `nn.Module`
 - We supplies two functions: `__init__()` and `forward()`
- Function $\rightarrow \dots \rightarrow$ Function \rightarrow Module/model definition (`class Net:...`) \rightarrow instantiation `net = Net(x)`
 - Instantiation wraps the internal `__call__()` method, so that we can do `y = net(x)`
- Any object which is of type “Variable” and is attached to the class definition will be automatically added for gradient computation
 - Unless we explicitly disable gradient computation

Generic enc-dec

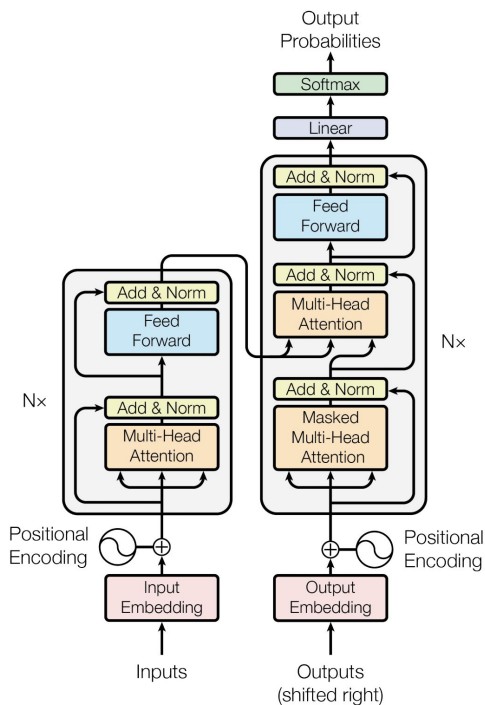


Figure 1: The Transformer - model architecture.

```
class EncoderDecoder(nn.Module):
```

```
    """
```

A standard Encoder-Decoder architecture. Base for this and many other models.

```
    """
```

```
    def __init__(self, encoder, decoder, src_embed, tgt_embed, generator):
```

```
        ...
```

```
    def forward(self, src, tgt, src_mask, tgt_mask):
```

"Take in and process masked src and target sequences."

```
        return self.decode(self.encode(src, src_mask), src_mask, tgt, tgt_mask)
```

```
    def encode(self, src, src_mask):
```

```
        return self.encoder(self.src_embed(src), src_mask)
```

```
    def decode(self, memory, src_mask, tgt, tgt_mask):
```

```
        return self.decoder(self.tgt_embed(tgt), memory, src_mask, tgt_mask)
```

Encoder

```
class EncoderLayer(nn.Module):  
    def __init__(self, size, self_attn, feed_forward, dropout):  
        super(EncoderLayer, self).__init__()  
        self.self_attn = self_attn  
        self.feed_forward = feed_forward  
        self.sublayer = clones(SublayerConnection(size, dropout), 2)  
        self.size = size  
  
    def forward(self, x, mask):  
        "Follow Figure 1 (left) for connections."  
        x = self.sublayer[0](x, lambda x: self.self_attn(x, x, x, mask))  
        return self.sublayer[1](x, self.feed_forward)
```

```
class Encoder(nn.Module):  
    "Core encoder is a stack of N layers"  
    def __init__(self, layer, N):  
        super(Encoder, self).__init__()  
        self.layers = clones(layer, N)  
        self.norm = LayerNorm(layer.size)  
  
    def forward(self, x, mask):  
        "Pass the input (and mask) through each layer in turn."  
        for layer in self.layers:  
            x = layer(x, mask)  
        return self.norm(x)
```

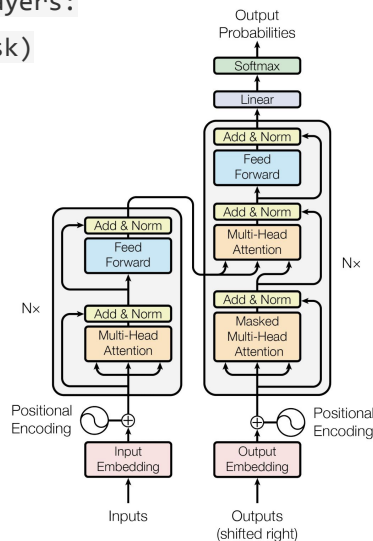


Figure 1: The Transformer - model architecture.

Encoder

```
class EncoderLayer(nn.Module):
```

```
    def __init__(self, size, self_attn, feed_forward, dropout):
```

```
        super(EncoderLayer, self).__init__()
```

```
        self.self_attn = self_attn
```

```
        self.feed_forward = feed_forward
```

```
        self.sublayer = clones(SublayerConnection(size, dropout), 2)
```

```
        self.size = size
```

```
    def forward(self, x, mask):
```

```
        "Follow Figure 1 (left) for connections."
```

```
        x = self.sublayer[0](x, lambda x: self.self_attn(x, x, x, mask))
```

```
        return self.sublayer[1](x, self.feed_forward)
```

```
class SublayerConnection(nn.Module):
```

```
    """
```

A residual connection followed by a layer norm.

Note for code simplicity the norm is first as opposed to last.

```
    """
```

```
    def __init__(self, size, dropout):
```

```
        super(SublayerConnection, self).__init__()
```

```
        self.norm = LayerNorm(size)
```

```
        self.dropout = nn.Dropout(dropout)
```

```
    def forward(self, x, sublayer):
```

```
        return x + self.dropout(sublayer(self.norm(x)))
```

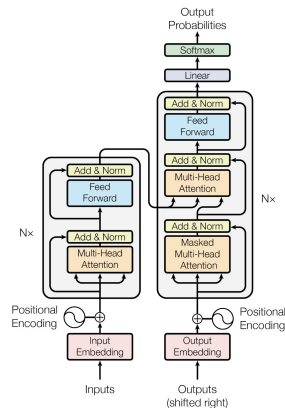


Figure 1: The Transformer - model architecture.

Decoder

```
class DecoderLayer(nn.Module):  
    def __init__(self, size, self_attn,  
                  src_attn, feed_forward, dropout):  
        super(DecoderLayer, self).__init__()  
        self.size = size  
        self.self_attn = self_attn  
        self.src_attn = src_attn  
        self.feed_forward = feed_forward  
        self.sublayer = clones(SublayerConnection(size, dropout), 3)
```

```
def forward(self, x, memory, src_mask, tgt_mask):  
    "Follow Figure 1 (right) for connections."  
    m = memory  
    x = self.sublayer[0](x, lambda x: self.self_attn(x, x, x, tgt_mask))  
    x = self.sublayer[1](x, lambda x: self.src_attn(x, m, m, src_mask))  
    return self.sublayer[2](x, self.feed_forward)
```

```
class Decoder(nn.Module):  
    "Generic N layer decoder with masking."  
    def __init__(self, layer, N):  
        super(Decoder, self).__init__()  
        self.layers = clones(layer, N)  
        self.norm = LayerNorm(layer.size)  
  
    def forward(self, x, memory, src_mask, tgt_mask):  
        for layer in self.layers:  
            x = layer(x, memory, src_mask, tgt_mask)  
        return self.norm(x)
```

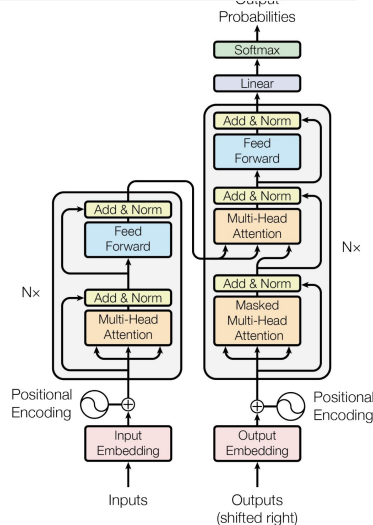


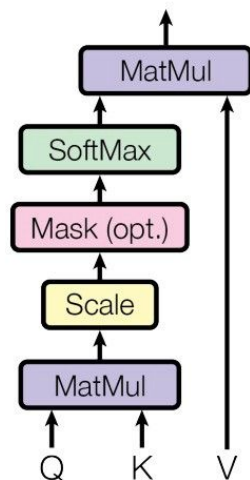
Figure 1: The Transformer - model architecture.

Attention overview

- **Keys:** A sequence of vectors also known as the memory
- **Values:** A sequence of vectors from which we aggregate the output through a weighted linear combination. Often Keys serve as Values.
- **Query:** A single vector that we use to probe the Keys
- **Output:** A single vector which is derived from a linear combination of the *Values* using the probabilities from the previous step as weights

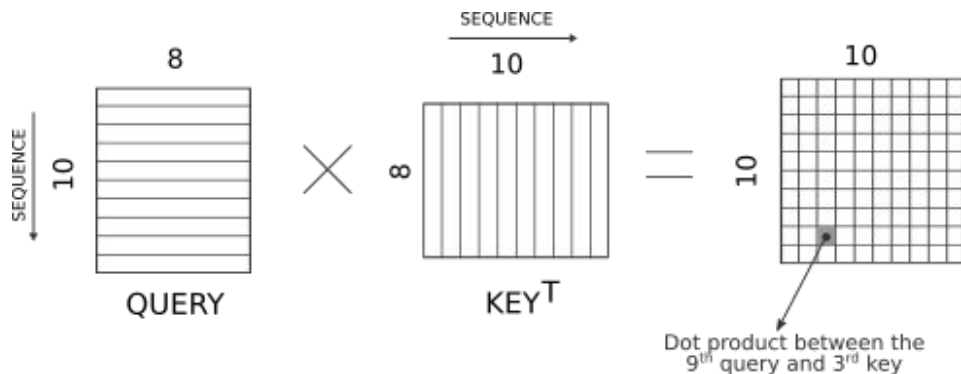
[Building the Mighty Transformer for Sequence Tagging in PyTorch](#)

Scaled Dot-Product Attention



$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K.$$



Attention, attention, attention

[Tensor2Tensor Transformers \(slides by Łukasz Kaiser\)](#)

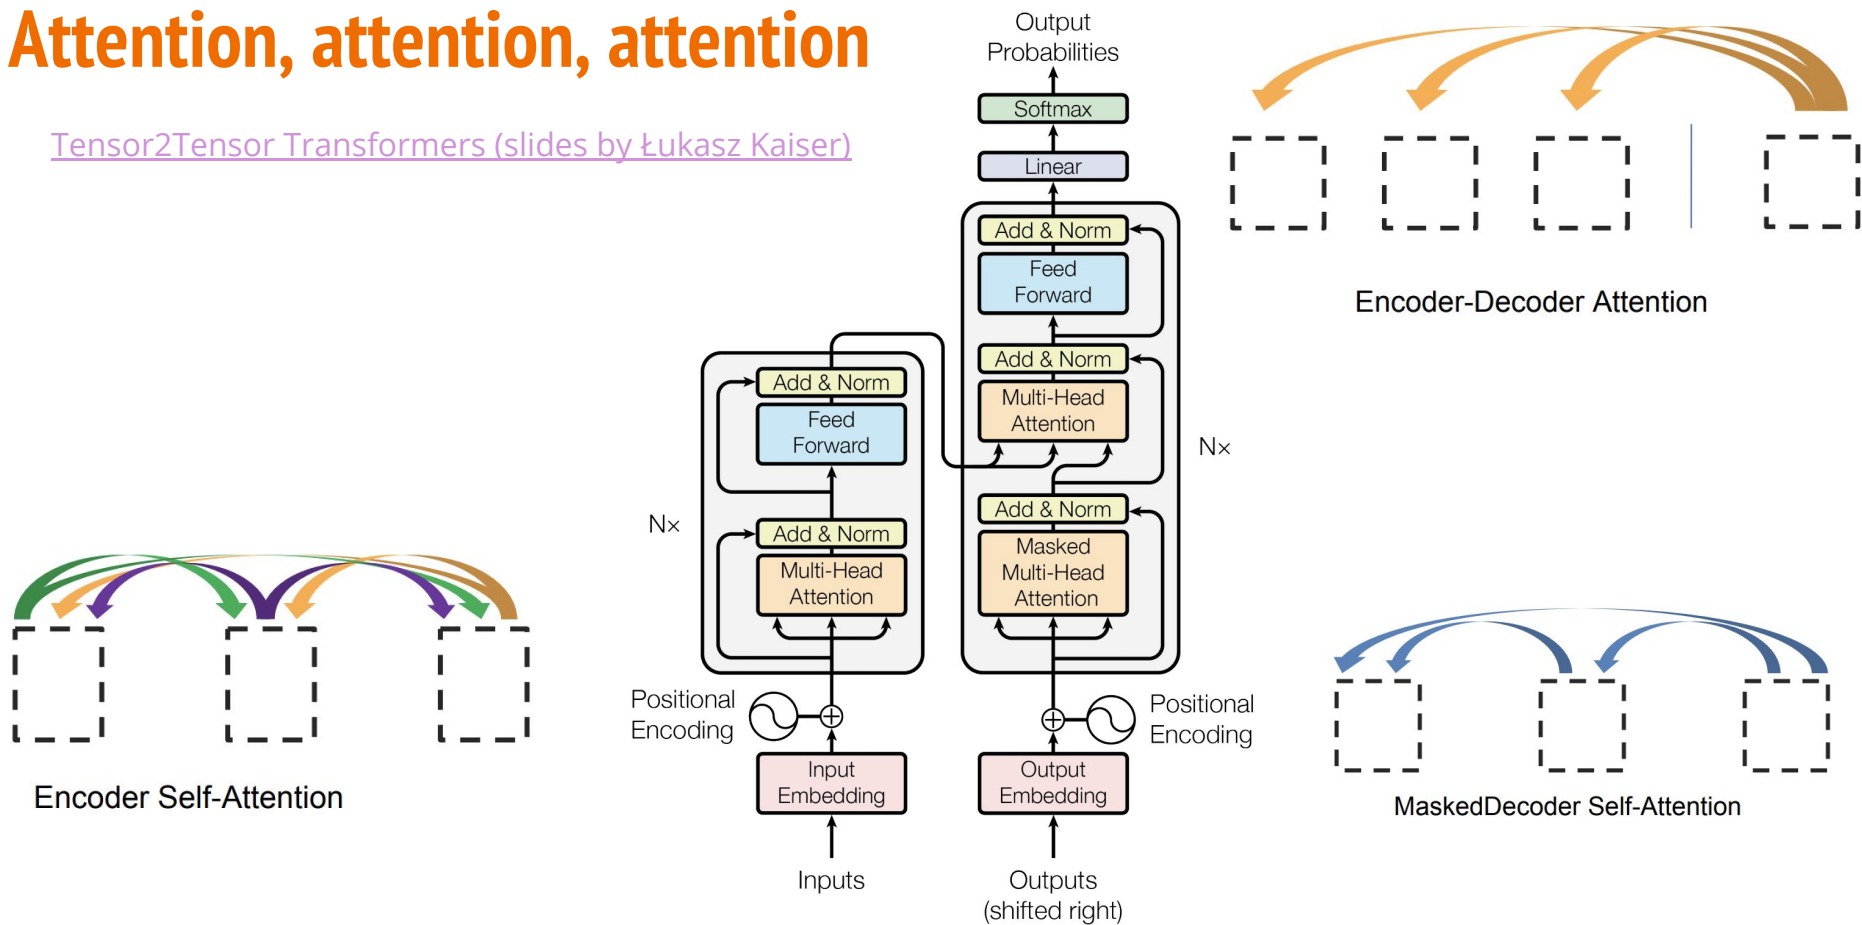
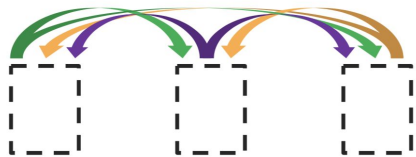


Figure 1: The Transformer - model architecture.

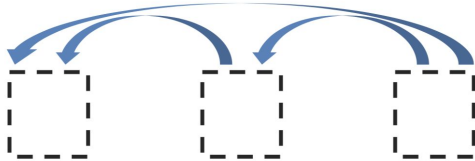
Self Attention vs enc-dec attention

```
self_attn_map = attention(x= $\text{bn} \times 8 \times 64 \times 512$ , x=ditto, x=ditto, mask)
```

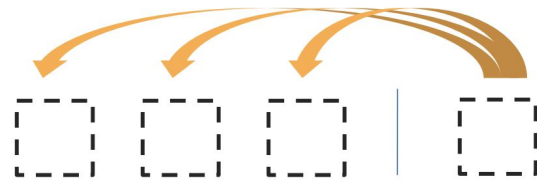
```
enc_dec_attn_map = attention(Q_split= $\text{bn} \times 8 \times 64 \times 512$ , K_split, V_split, mask)
```



Encoder Self-Attention



MaskedDecoder Self-Attention



Encoder-Decoder Attention

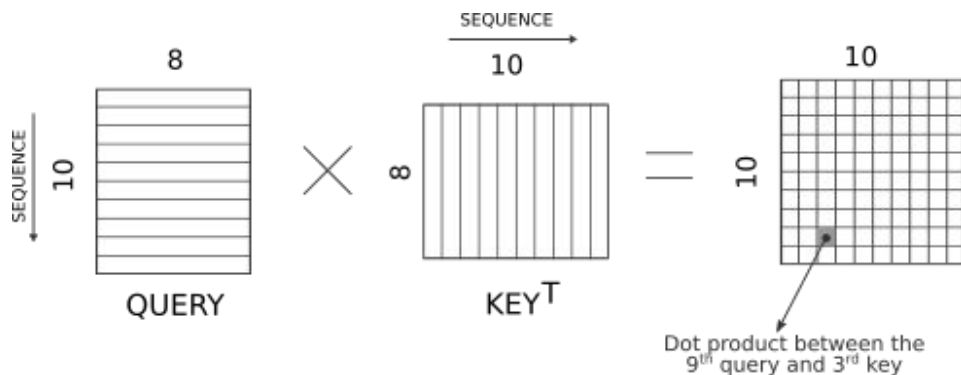
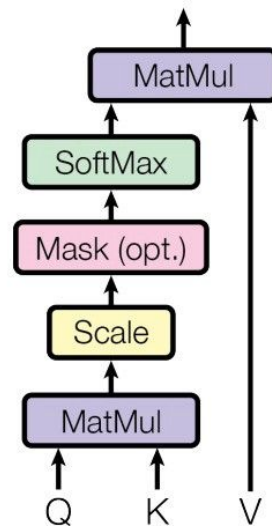
Generic Attention

```
def attention(querybn*H*S*D, keybn*H*S*D, valuebn*H*S*D, mask = None):
    d_kD = query.size(-1)
    scores = torch.matmul(query, key.transpose(-2, -1)) / math.sqrt(d_k)
    if mask is not None:
        scores = scores.masked_fill(mask == 0, -1e9)
    p_attnbn*H*S*S = F.softmax(scores, dim = -1)
    return torch.matmul(p_attnbn*H*S*S, valuebn*H*S*D)bn*H*S*D, p_attn
```

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K.$$

Q, K & V:4D Tensors
[batch_size, num_heads, seq_len, depth/num_heads]



Multi-Head Attention

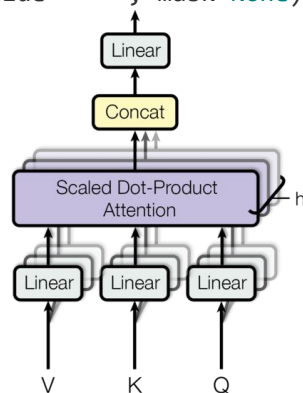
```
class MultiHeadedAttention(nn.Module):
    def __init__(self, h=8, d_model=512, dropout=0.1):
        "Take in model size and number of heads."
        super(MultiHeadedAttention, self).__init__()
        assert d_model % h == 0
        # We assume d_v always equals d_k
        self.d_k = 64 = d_model // h
        self.h = h
        self.linears = clones(
            nn.Linear(d_model, d_model)=512*512, 4
        )
        self.attn = None
        self.dropout = nn.Dropout(p=dropout)
```

```
def forward(self, query=bs*512*512, key=ditto, value=ditto, mask=None):
    "Implements Figure 2"
    if mask is not None:
        # Same mask applied to all h heads.
        mask = mask.unsqueeze(1)
    nbatches = query.size(0)

    # 1) Do all the linear projections in batch
    # from d_model => h x d_k
    query=bs*512*8*64, key=ditto, value=ditto = \
        [l(x).view(nbatches, -1, self.h, self.d_k).transpose(1, 2)
         for l, x in zip(self.linears, (query, key, value))]

    # 2) Apply attention on all the projected vectors in batch.
    x, self.attn = attention(query, key, value, mask=mask,
                             dropout=self.dropout)

    # 3) "Concat" using a view and apply a final linear.
    x = x.transpose(1, 2).contiguous() \
        .view(nbatches, -1, self.h * self.d_k)
    return self.linears[-1](x)
```



Masking

```
def _gen_bias_mask(max_length):
```

```
    """
```

```
    Generates bias values (-Inf) to mask future timesteps during attention
```

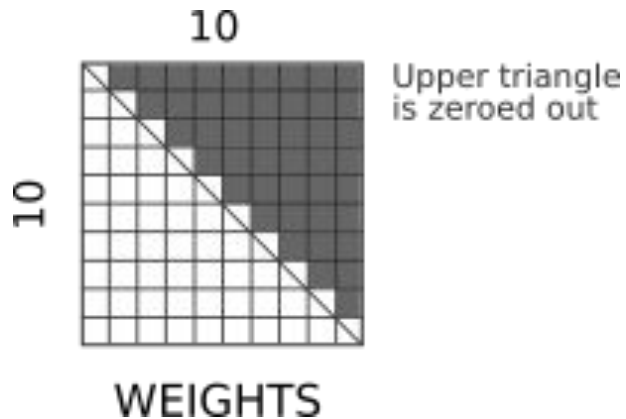
```
    """
```

```
    np_mask = np.triu(np.full([max_length, max_length], -np.inf), 1)
```

```
    torch_mask = torch.from_numpy(np_mask).type(torch.FloatTensor)
```

```
# Reshape to 4D Tensor to handle multiple heads
```

```
return torch_mask.unsqueeze(0).unsqueeze(1)
```



```
>>> np.triu([[1,2,3],[4,5,6],[7,8,9],[10,11,12]], 1)
array([[ 0,  2,  3],
       [ 0,  0,  6],
       [ 0,  0,  9],
       [ 0,  0,  0]])
```

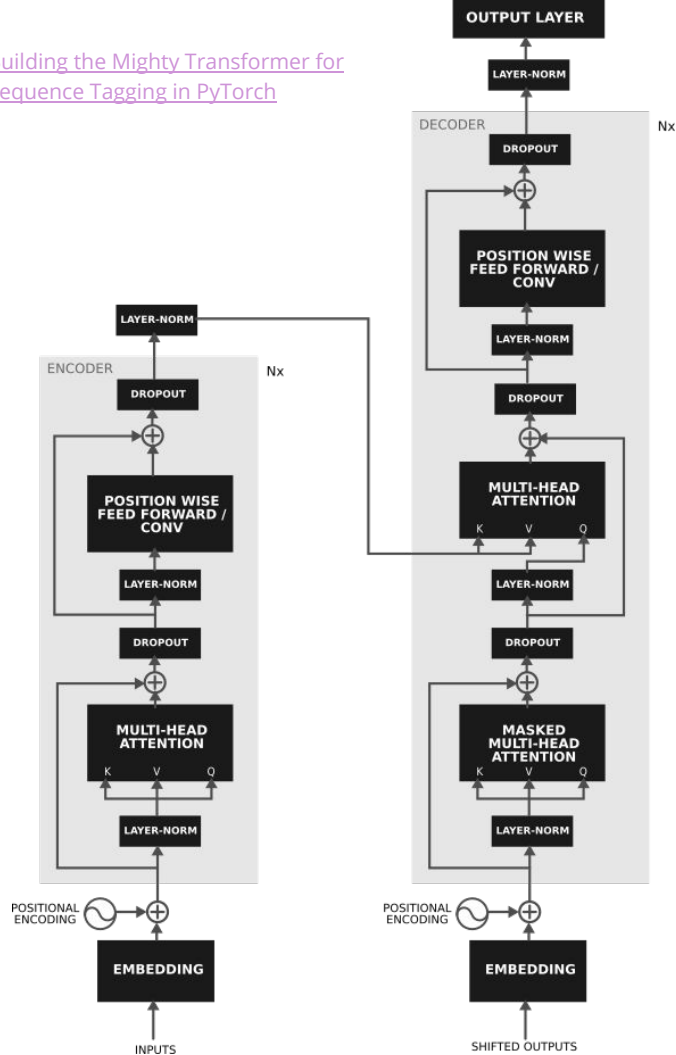
Positionwise Feed Forward

- applied to each position *separately* and *identically*

```
class PositionwiseFeedForward(nn.Module):  
    "Implements FFN equation."  
  
    def __init__(self, d_model=512, d_ff=2048, dropout=0.1):  
        super(PositionwiseFeedForward, self).__init__()  
        self.w_1=512x2048 = nn.Linear(d_model, d_ff)  
        self.w_2=2048x512 = nn.Linear(d_ff, d_model)  
        self.dropout = nn.Dropout(dropout)  
  
    def forward=bn*S*512(self, x=bn*S*512):  
        return self.w_2(self.dropout(F.relu(self.w_1(x))))
```

$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$

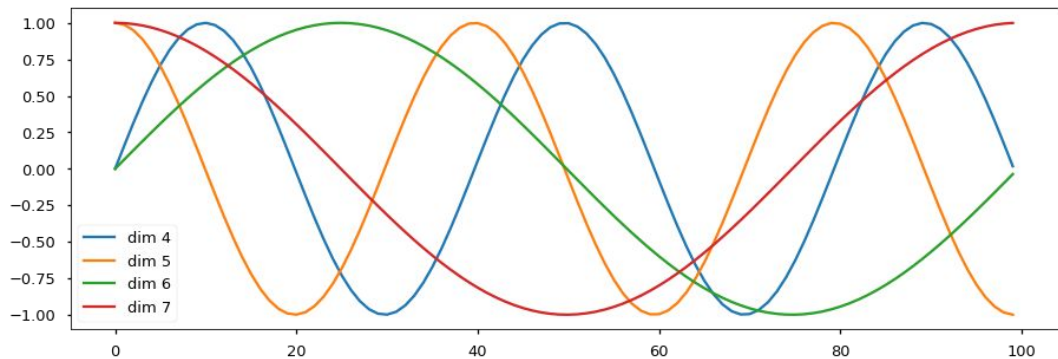
[Building the Mighty Transformer for Sequence Tagging in PyTorch](#)



Positional Encoding

Gehring, Jonas, et al. "Convolutional sequence to sequence learning."
arXiv preprint arXiv:1705.03122 (2017).

```
class PositionalEncoding(nn.Module):  
    def __init__(self, d_model=512, dropout, max_len=5000):  
        ...  
        self.dropout = nn.Dropout(p=dropout)  
        # Compute the positional encodings once in log space.  
        pe = torch.zeros(max_len, d_model)  
        position = torch.arange(0, max_len).unsqueeze(1)  
        div_term = torch.exp(torch.arange(0, d_model, 2) *  
                               -(math.log(10000.0) / d_model))  
        pe[:, 0::2] = torch.sin(position * div_term)  
        pe[:, 1::2] = torch.cos(position * div_term)  
        pe = pe.unsqueeze(0)  
        self.register_buffer('pe', pe)  
  
    def forward(self, x):  
        x = x + Variable(self.pe[:, :x.size(1)],  
                          requires_grad=False)  
        return self.dropout(x)
```



LayerNorm

```
class LayerNorm(nn.Module):  
    "Construct a layernorm module (See citation for details)."  
    def __init__(self, features, eps=1e-6):  
        super(LayerNorm, self).__init__()  
        self.a_2 = nn.Parameter(torch.ones(features))  
        self.b_2 = nn.Parameter(torch.zeros(features))  
        self.eps = eps  
  
    def forward(self, x):  
        mean = x.mean(-1, keepdim=True)  
        std = x.std(-1, keepdim=True)  
        return self.a_2 * (x - mean) / (std + self.eps) + self.b_2
```

Put Together

```
def make_model(src_vocab, tgt_vocab, N=6,
               d_model=512, d_ff=2048, h=8, dropout=0.1):
    "Helper: Construct a model from hyperparameters."
    c = copy.deepcopy
    attn = MultiHeadedAttention(h, d_model)
    ff = PositionwiseFeedForward(d_model, d_ff, dropout)
    position = PositionalEncoding(d_model, dropout)
    model = EncoderDecoder(
        Encoder(EncoderLayer(d_model, c(attn), c(ff), dropout), N),
        Decoder(DecoderLayer(d_model, c(attn), c(attn),
                              c(ff), dropout), N),
        nn.Sequential(Embeddings(d_model, src_vocab), c(position)),
        nn.Sequential(Embeddings(d_model, tgt_vocab), c(position)),
        Generator(d_model, tgt_vocab))
    # This was important from their code.
    # Initialize parameters with Glorot / fan_avg.
    for p in model.parameters():
        if p.dim() > 1:
            nn.init.xavier_uniform(p)
    return model
```

Train it

```
class SimpleLossCompute:
```

```
    "A simple loss compute and train function."
```

```
    def __init__(self, generator, criterion, optimizer=None):
        ...
```

```
    def __call__(self, x, y, norm):
        x = self.generator(x)
        loss = self.criterion(x.contiguous().view(-1, x.size(-1)),
                               y.contiguous().view(-1)) / norm

        loss.backward()

        if self.opt is not None:
            self.opt.step()
            self.opt.optimizer.zero_grad()

        return loss.data[0] * norm
```

```
def run_epoch(data_iter, model, loss_compute):
    "Standard Training and Logging Function"
    start = time.time()
    total_tokens = 0
    total_loss = 0
    tokens = 0
    for i, batch in enumerate(data_iter):
        out = model.forward(batch.src, batch.trg,
                             batch.src_mask, batch.trg_mask)
        loss = loss_compute(out, batch.trg_y, batch.ntokens)
        total_loss += loss
        total_tokens += batch.ntokens
        tokens += batch.ntokens
        if i % 50 == 1:
            elapsed = time.time() - start
            print("Epoch Step: %d Loss: %f Tokens per Sec: %f" %
                  (i, loss / batch.ntokens, tokens / elapsed))
            start = time.time()
            tokens = 0
    return total_loss / total_tokens
```

Train it

```
def run_epoch(data_iter, model, loss_compute):
    "Standard Training and Logging Function"
    start = time.time()
    total_tokens = 0
    total_loss = 0
    tokens = 0
    for i, batch in enumerate(data_iter):
        out = model.forward(batch.src, batch.trg,
                             batch.src_mask, batch.trg_mask)
        loss = loss_compute(out, batch.trg_y, batch.ntokens)
        total_loss += loss
        total_tokens += batch.ntokens
        tokens += batch.ntokens
        if i % 50 == 1:
            elapsed = time.time() - start
            print("Epoch Step: %d Loss: %f Tokens per Sec: %f" %
                  (i, loss / batch.ntokens, tokens / elapsed))
            start = time.time()
            tokens = 0
    return total_loss / total_tokens
```

Epoch Step: 1 Loss: 3.023465 Tokens per Sec: 403.074173

Epoch Step: 1 Loss: 1.920030 Tokens per Sec: 641.689380

1.9274832487106324

Epoch Step: 1 Loss: 1.940011 Tokens per Sec: 432.003378

Epoch Step: 1 Loss: 1.699767 Tokens per Sec: 641.979665

1.657595729827881

Epoch Step: 1 Loss: 1.860276 Tokens per Sec: 433.320240

Epoch Step: 1 Loss: 1.546011 Tokens per Sec: 640.537198

1.4888023376464843

...

Epoch Step: 1 Loss: 0.459483 Tokens per Sec: 434.594030

Epoch Step: 1 Loss: 0.290385 Tokens per Sec: 642.519464

0.2612409472465515

Epoch Step: 1 Loss: 1.031042 Tokens per Sec: 434.557008

Epoch Step: 1 Loss: 0.437069 Tokens per Sec: 643.630322

0.4323212027549744

Epoch Step: 1 Loss: 0.617165 Tokens per Sec: 436.652626

Epoch Step: 1 Loss: 0.258793 Tokens per Sec: 644.372296

0.27331129014492034

Optimizer

We used the Adam optimizer (cite) with $\beta_1 = 0.9$, $\beta_2 = 0.98$ and $\epsilon = 10^{-9}$. We varied the learning rate over the course of training, according to the formula: $lrate = d_{\text{model}}^{-0.5} \cdot \min(\text{step_num}^{-0.5}, \text{step_num} \cdot \text{warmup_steps}^{-1.5})$ This corresponds to increasing the learning rate linearly for the first warmup_steps training steps, and decreasing it thereafter proportionally to the inverse square root of the step number. We used $\text{warmup_steps} = 4000$.

```
class NoamOpt:
```

```
    "Optim wrapper that implements rate."
```

```
    def __init__(self, model_size, factor, warmup, optimizer):
```

```
        ...
```

```
        self._step = 0
```

```
        self._rate = 0
```

```
    def step(self):
```

```
        "Update parameters and rate"
```

```
        self._step += 1
```

```
        rate = self.rate()
```

```
        for p in self.optimizer.param_groups:
```

```
            p['lr'] = rate
```

```
        self._rate = rate
```

```
        self.optimizer.step()
```

```
    "Implement `lrate` above"
```

```
    if step is None:
```

```
        step = self._step
```

```
    return self.factor * \
```

```
        (self.model_size ** (-0.5)) *
```

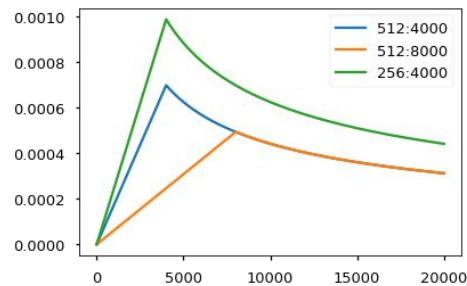
```
        min(step ** (-0.5), step * self.warmup ** (-1.5)))
```

```
def get_std_opt(model):
```

```
    return NoamOpt(model.src_embed[0].d_model, 2, 4000,
```

```
                    torch.optim.Adam(model.parameters(), lr=0,
```

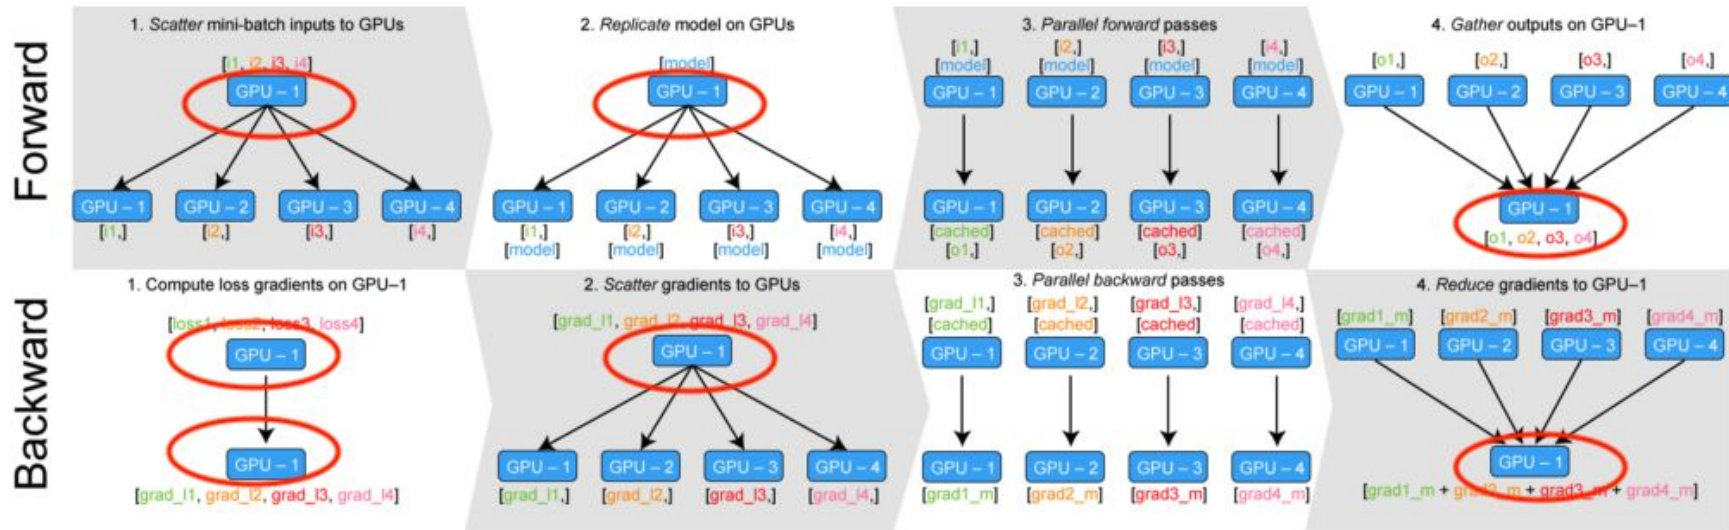
```
                    betas=(0.9, 0.98), eps=1e-9))
```



$$lrate = d_{\text{model}}^{-0.5} \cdot \min(\text{step_num}^{-0.5}, \text{step_num} \cdot \text{warmup_steps}^{-1.5})$$

Multi-GPU training

- replicate - split modules onto different gpus.
- scatter - split batches onto different gpus
- parallel_apply - apply module to batches on different gpus
- gather - pull scattered data back onto one gpu.
- nn.DataParallel - a special module wrapper that calls these all before evaluating.



Forward and Backward passes with torch.nn.DataParallel

```

class MultiGPULossCompute:
    ...
    def __call__(self, out, targets, normalize):
        total = 0.0
        generator = nn.parallel.replicate(self.generator,
            devices=self.devices)
        out_scatter = nn.parallel.scatter(out, target_gpus=self.devices)
        out_grad = [[] for _ in out_scatter]
        targets = nn.parallel.scatter(targets, target_gpus=self.devices)
        # Divide generating into chunks.
        chunk_size = self.chunk_size
        for i in range(0, out_scatter[0].size(1), chunk_size):
            # Predict distributions
            out_column = [[Variable(o[:, i:i+chunk_size].data,
                requires_grad=self.opt is not None)]
                for o in out_scatter]
            gen = nn.parallel.parallel_apply(generator, out_column)
            # Compute Loss.
            y = [(g.contiguous().view(-1, g.size(-1)),
                t[:, i:i+chunk_size].contiguous().view(-1))
                for g, t in zip(gen, targets)]
            loss = nn.parallel.parallel_apply(self.criterion, y)

```

```

        # Sum and normalize loss
        l = nn.parallel.gather(loss,
            target_device=self.devices[0])
        l = l.sum()[0] / normalize
        total += l.data[0]
        # Backprop loss to output of transformer
        if self.opt is not None:
            l.backward()
            for j, l in enumerate(loss):
                out_grad[j].append(out_column[j][0].grad.data.clone())
                # Backprop all loss through transformer.
                if self.opt is not None:
                    out_grad = [Variable(torch.cat(og, dim=1))
                        for og in out_grad]
                    o1 = out
                    o2 = nn.parallel.gather(out_grad,
                        target_device=self.devices[0])
                    o1.backward(gradient=o2)
                    self.opt.step()
                    self.opt.optimizer.zero_grad()
            return total * normalize

```

Test it out: greedy decoding

```
def greedy_decode(model, src, src_mask, max_len, start_symbol):  
    memory = model.encode(src, src_mask)  
    ys = torch.ones(1, 1).fill_(start_symbol).type_as(src.data)  
    for i in range(max_len-1):  
        out = model.decode(memory, src_mask,  
                           Variable(ys),  
                           Variable(subsequent_mask(ys.size(1))  
                                   .type_as(src.data)))  
        prob = model.generator(out[:, -1])  
        _, next_word = torch.max(prob, dim = 1)  
        next_word = next_word.item()  
        ys = torch.cat([  
            ys, torch.ones(1, 1)  
                .type_as(src.data).fill_(next_word)], dim=1)  
    return ys
```

```
model.eval()  
src = Variable(torch.LongTensor([[1,2,3,4,5,6,7,8,9,10]]))  
src_mask = Variable(torch.ones(1, 1, 10))  
print(greedy_decode(model, src, src_mask, max_len=10,  
start_symbol=1))
```

```
1 2 3 4 5 6 7 8 9 10  
[torch.LongTensor of size 1x10]
```

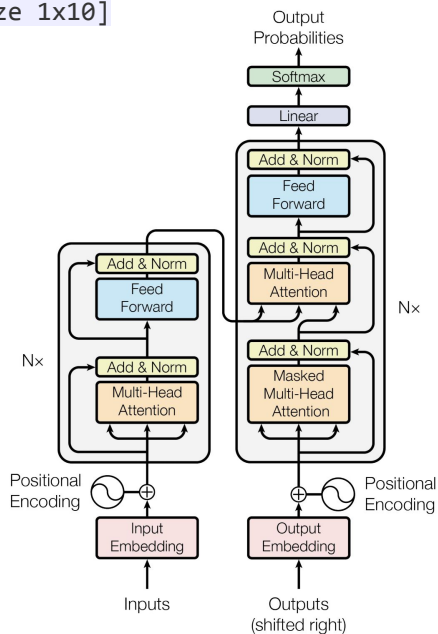


Figure 1: The Transformer - model architecture.

The Authors' Implementation

- <https://github.com/tensorflow/tensor2tensor/blob/master/tensor2tensor/models/transformer.py>
- TF Estimator
- Supports TPU, GPU & CPU
- Attention caching

Thanks

- Thanks to Renyu Li and another friend for helping with the preparation
- Thanks to Rouzbeh Afrasiabi and Dave Fernandes for suggesions

References

- Vaswani, Ashish, et al. "[Attention is all you need.](#)" Advances in Neural Information Processing Systems. 2017
- [The Annotated Transformer](#)
- [The Illustrated Transformer](#)
- [How to code The Transformer in Pytorch](#)
- [Michał Chromiak's blog post on Transformer](#)
- [Tensor2Tensor Transformers \(slides by Łukasz Kaiser\)](#)
- [Building the Mighty Transformer for Sequence Tagging in PyTorch](#)
- [Transformer from NLP Tutorial by Tae Hwan Jung\(Jeff Jung\) & Derek Miller](#)

Model averaging

The paper averages the last k checkpoints to create an ensembling effect.

```
def average(model, models):  
    "Average models into model"  
    for ps in zip(*[m.params() for m in [model] + models]):  
        p[0].copy_(torch.sum(*ps[1:]) / len(ps[1:]))
```